

# Project 1

## External Sorting

# External Sorting

- Sort  $n$  records/elements that reside on a disk.
- Space needed by the  $n$  records is very large.
  - $n$  is very large, and each record may be large or small.
  - $n$  is small, but each record is very large.
- So, not feasible to input the  $n$  records, sort, and output in sorted order.


# Small $n$ But Large File

- Input the record keys.
- Sort the  $n$  keys to determine the sorted order for the  $n$  records.
- Permute the records into the desired order (possibly several fields at a time).
- We focus on the case: large  $n$ , large file.


# Data Structures/Concepts

- Tournament trees.
- Huffman trees.
- Double-ended priority queues.
- Buffering.
- Ideas also may be used to speed algorithms for small instances by using cache more efficiently.

# External Sort Computer Model




# Disk Characteristics




- Seek time
  - Approx. 100,000 arithmetics
- Latency time
  - Approx. 25,000 arithmetics
- Transfer time
- Data access by block

# Traditional Internal Memory Model



# More Accurate Memory Model





# Phase 1

Warm up


# Matrix Multiplication

```
for (int i = 0; i < n; i++)  
    for (int j = 0; j < n; j++)  
        for (int k = 0; k < n; k++)  
            c[i][j] += a[i][k] * b[k][j];
```

- ijk, ikj, jik, jki, kij, kji orders of loops yield same result.
- All perform same number of operations.
- But run time may differ significantly!

# 2D Array Representation In Java, C, and C++


```
int x[3][4];
```




## Array of Arrays Representation

# ijk Order

```
for (int i = 0; i < n; i++)  
  for (int j = 0; j < n; j++)  
    for (int k = 0; k < n; k++)  
      c[i][j] += a[i][k] * b[k][j];
```




# ijk Analysis




- Block size = width of cache line =  $w$ .
- Assume one-level cache.
- $C \Rightarrow n^2/w$  cache misses.
- $A \Rightarrow n^3/w$  cache misses, when  $n$  is large.
- $B \Rightarrow n^3$  cache misses, when  $n$  is large.
- Total cache misses =  $n^3/w(1/n + 1 + w)$ .

# ikj Order

```
for (int i = 0; i < n; i++)  
  for (int k = 0; k < n; k++)  
    for (int j = 0; j < n; j++)  
      c[i][j] += a[i][k] * b[k][j];
```



# ikj Analysis



- $C \Rightarrow n^3/w$  cache misses, when  $n$  is large.
- $A \Rightarrow n^2/w$  cache misses.
- $B \Rightarrow n^3/w$  cache misses, when  $n$  is large.
- Total cache misses =  $n^3/w(2 + 1/n)$ .

# Warm up project

- Cache miss simulation
  - Files
- Count the missed hits
- #mh vs cache line
- #mh vs data volumn
- Simulation vs Theoretical results
- Others ...
- <https://www.cplusplus.com>



# Phase 2

External Sort: Quick Sort

# External Sort Methods

- Base the external sort method on a fast internal sort method.
- Average run time
  - Quick sort
- Worst-case run time
  - Merge sort

# Internal Quick Sort

- To sort a large instance, select a **pivot** element from out of the **n** elements.
- Partition the **n** elements into **3** groups **left**, **middle** and **right**.
- The **middle** group contains only the **pivot** element.
- All elements in the **left** group are  $\leq$  **pivot**.
- All elements in the **right** group are  $\geq$  **pivot**.
- Sort **left** and **right** groups recursively.
- Answer is sorted **left** group, followed by **middle** group followed by sorted **right** group.

# Internal Quick Sort


6	2	8	5	11	10	4	1	9	7	3
---	---	---	---	----	----	---	---	---	---	---

Use 6 as the pivot.

2	5	4	1	3	6	7	9	10	11	8
---	---	---	---	---	---	---	---	----	----	---


Sort left and right groups recursively.

# Quick Sort – External Adaptation




- 3 input/output buffers
  - input, small, large
- rest is used for middle group

# Quick Sort – External Adaptation



- fill middle group from disk
- if next **record**  $\leq$   $\text{middle}_{\min}$  send to **small**
- if next **record**  $\geq$   $\text{middle}_{\max}$  send to **large**
- **else** remove  $\text{middle}_{\min}$  or  $\text{middle}_{\max}$  from **middle** and add new record to middle group

# Quick Sort – External Adaptation




- Fill **input** buffer when it gets empty.
- Write **small/large** buffer when full.
- Write **middle** group in sorted order when done.
- Double-ended priority queue.

# Double-ended Priority Queue: Interval Heaps

- Complete binary tree.
- Each node (except possibly last one) has 2 elements.
- Last node has 1 or 2 elements.
- Let  $a$  and  $b$  be the elements in a node  $P$ ,  $a \leq b$ .
- $[a, b]$  is the interval represented by  $P$ .
- The interval represented by a node that has just one element  $a$  is  $[a, a]$ .
- The interval  $[c, d]$  is contained in interval  $[a, b]$  iff  $a \leq c \leq d \leq b$ .
- In an interval heap each node's (except for root) interval is contained in that of its parent.




# Interval

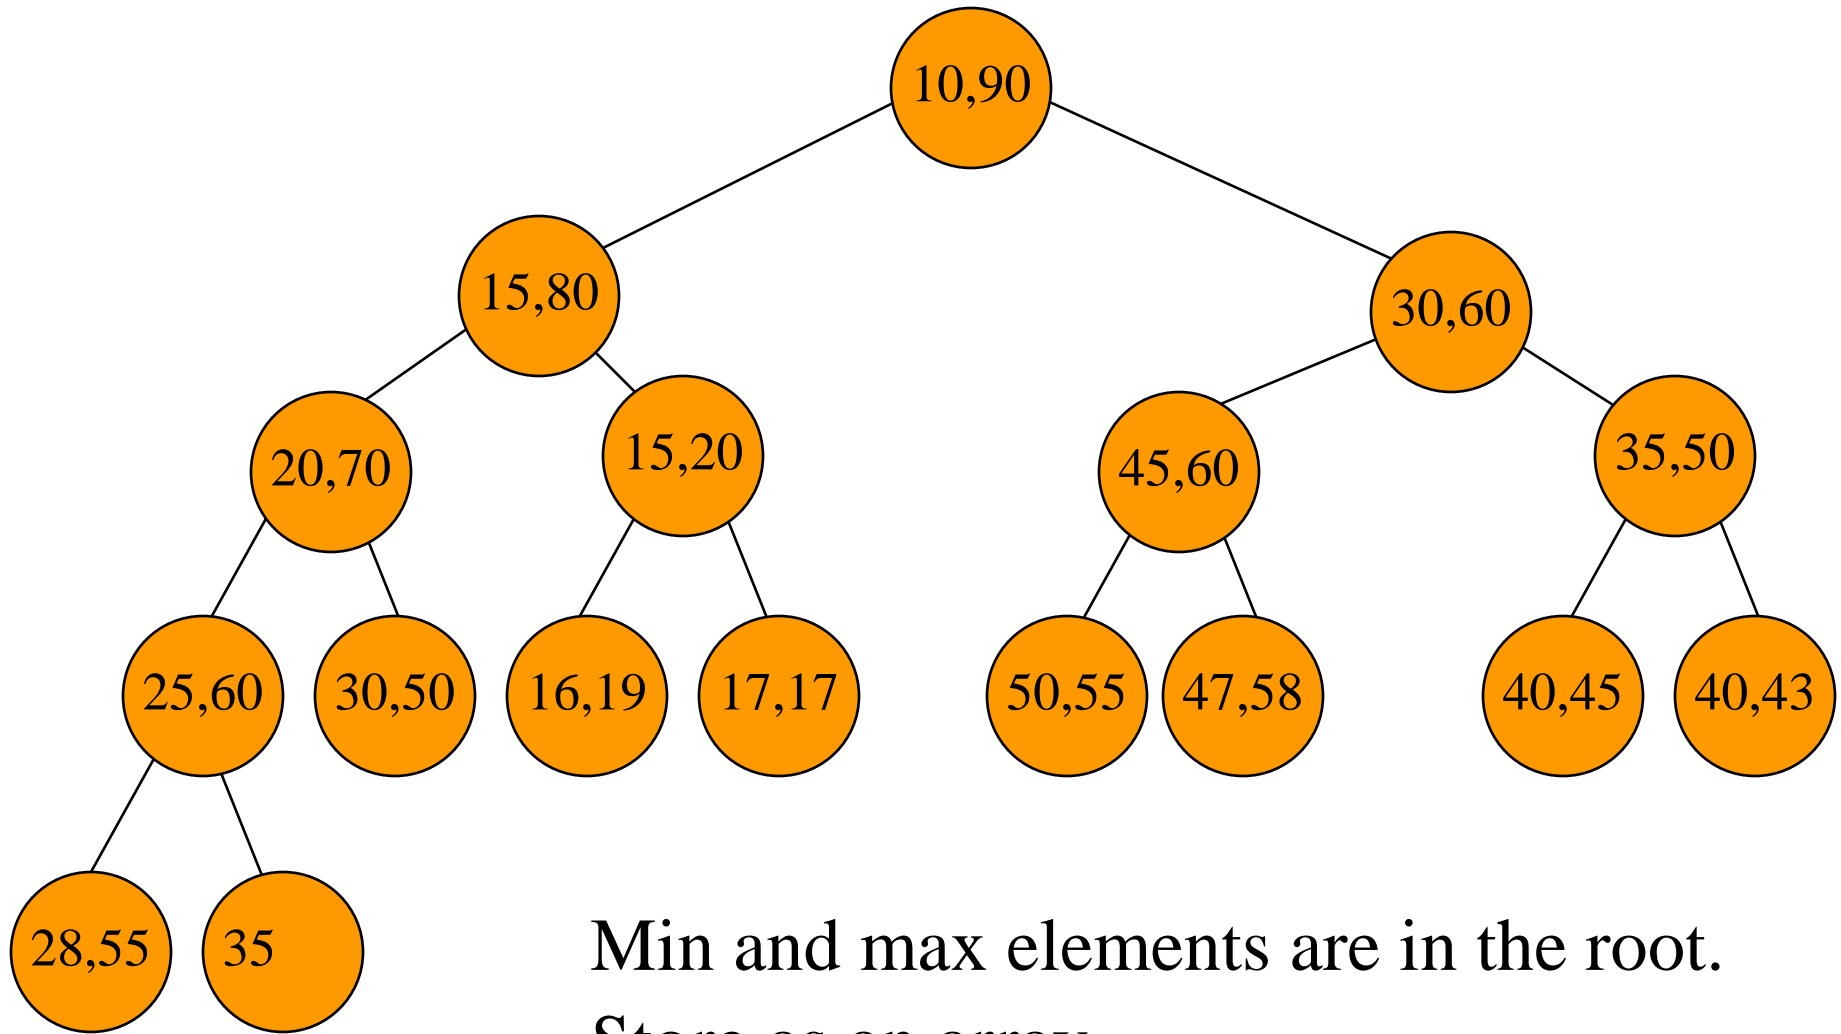


- $[c, d]$  is contained in  $[a, b]$
- $a \leq c$
- $d \leq b$

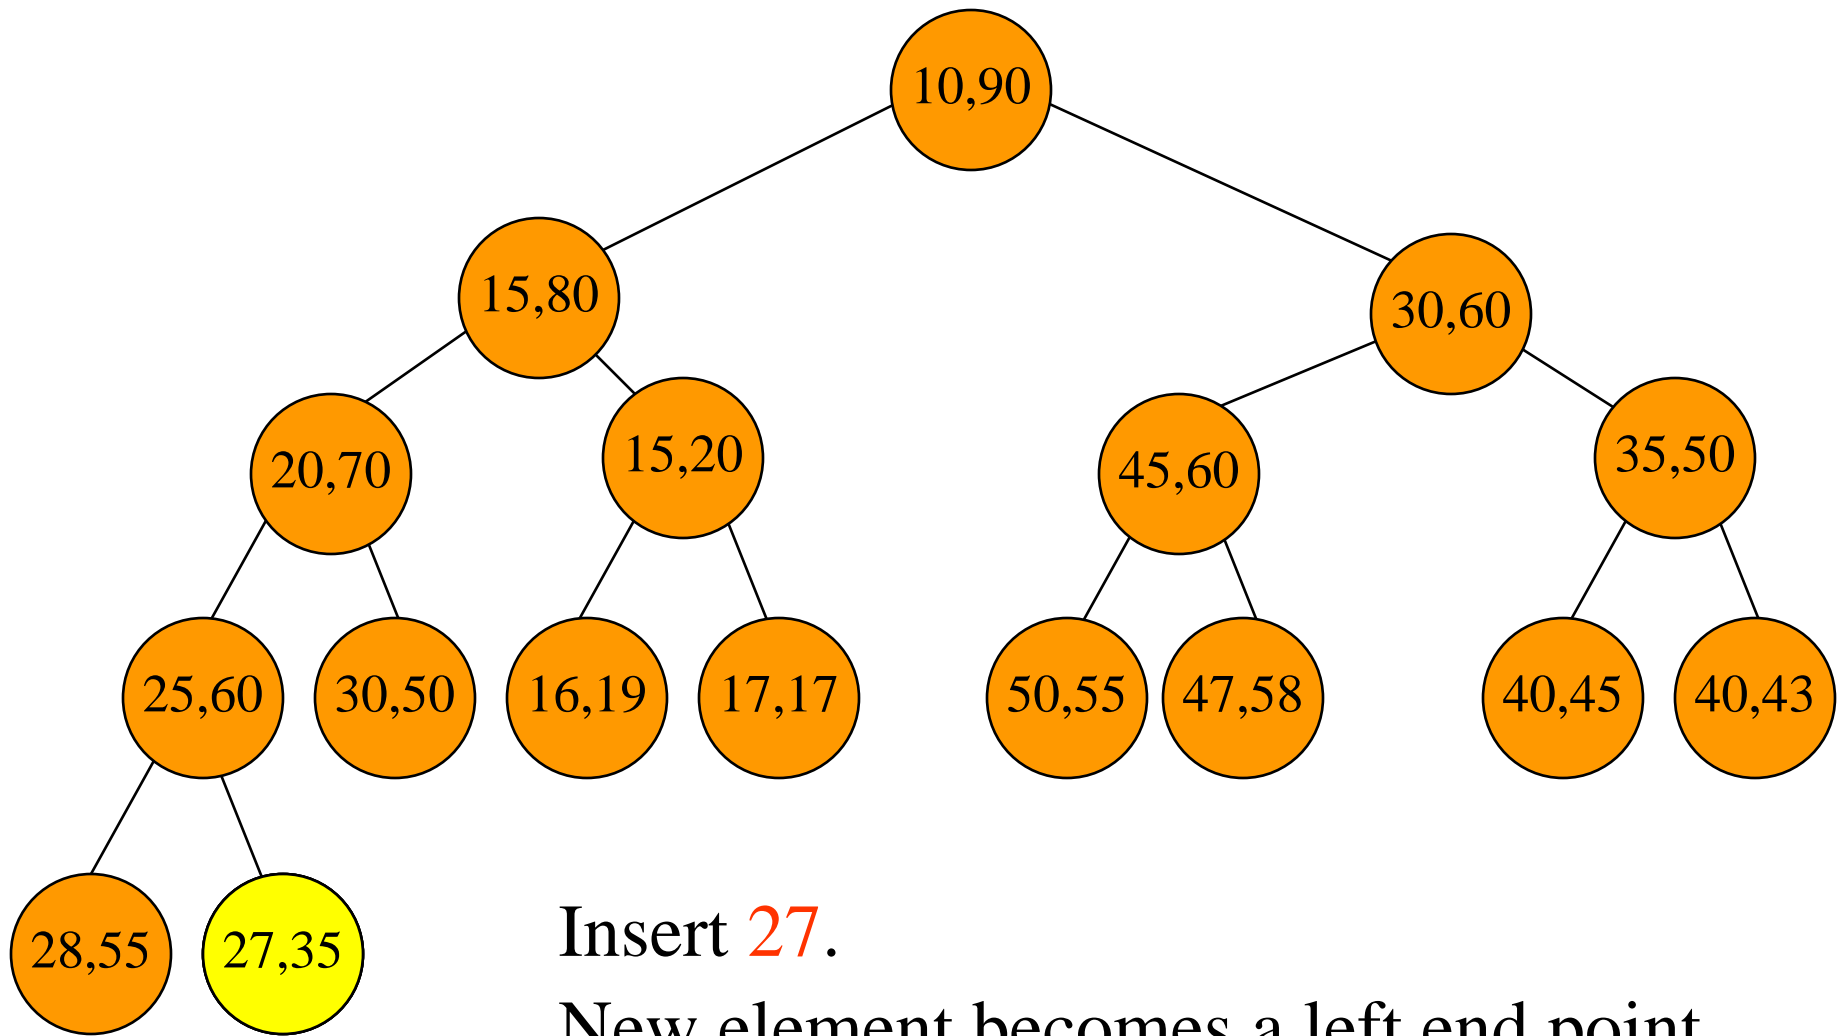
# Example Interval Heap



# Example Interval Heap



# Insert An Element

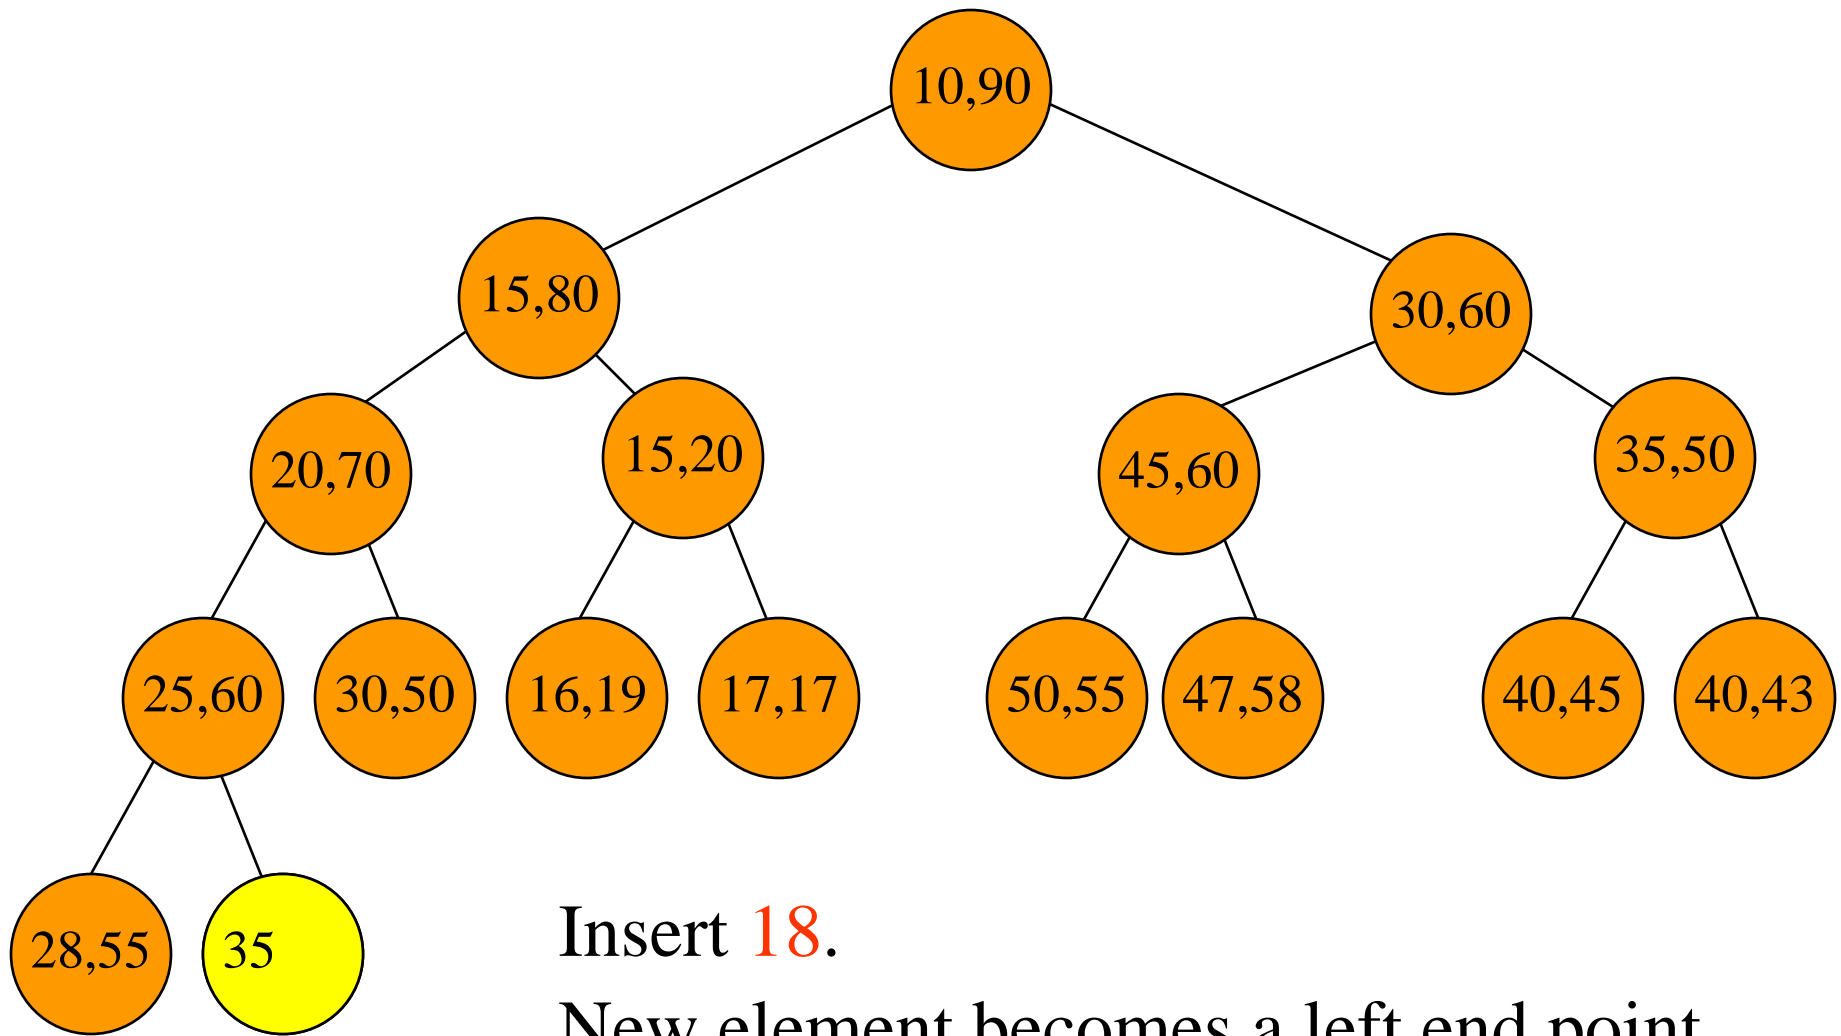


Insert **27**.

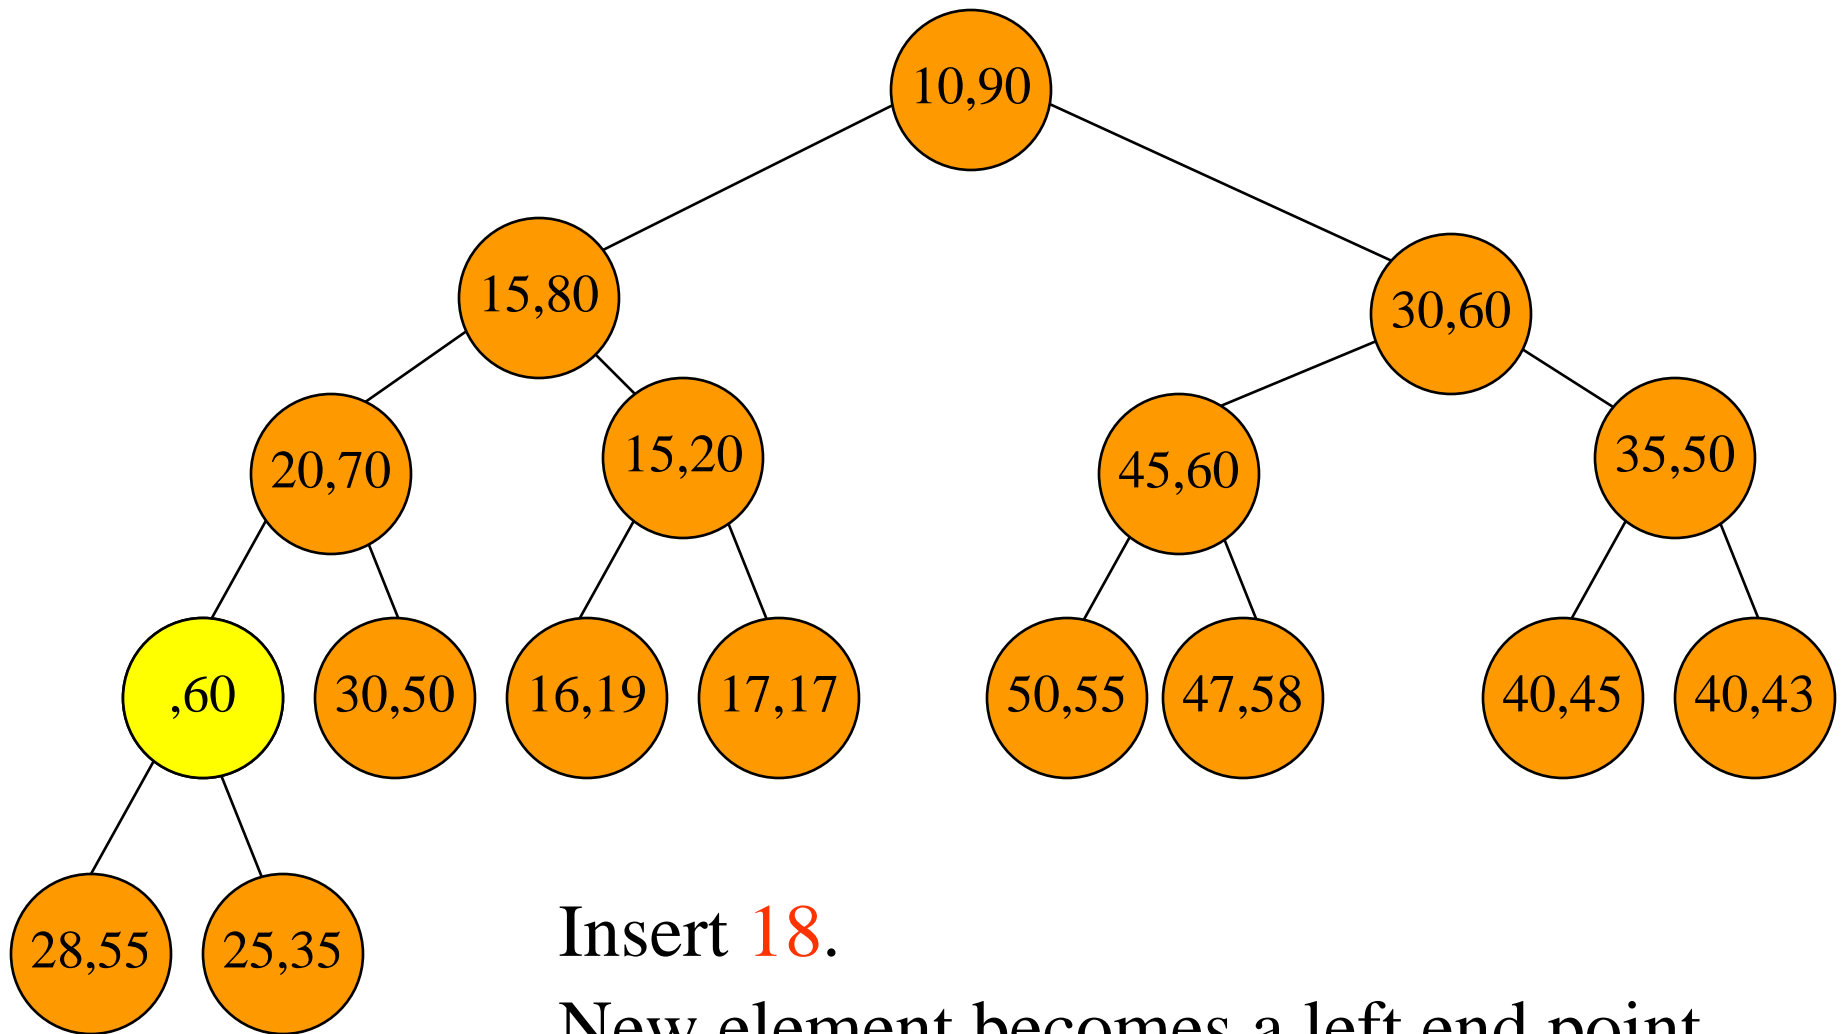
New element becomes a left end point.

Insert new element into min heap.

# Another Insert



# Another Insert

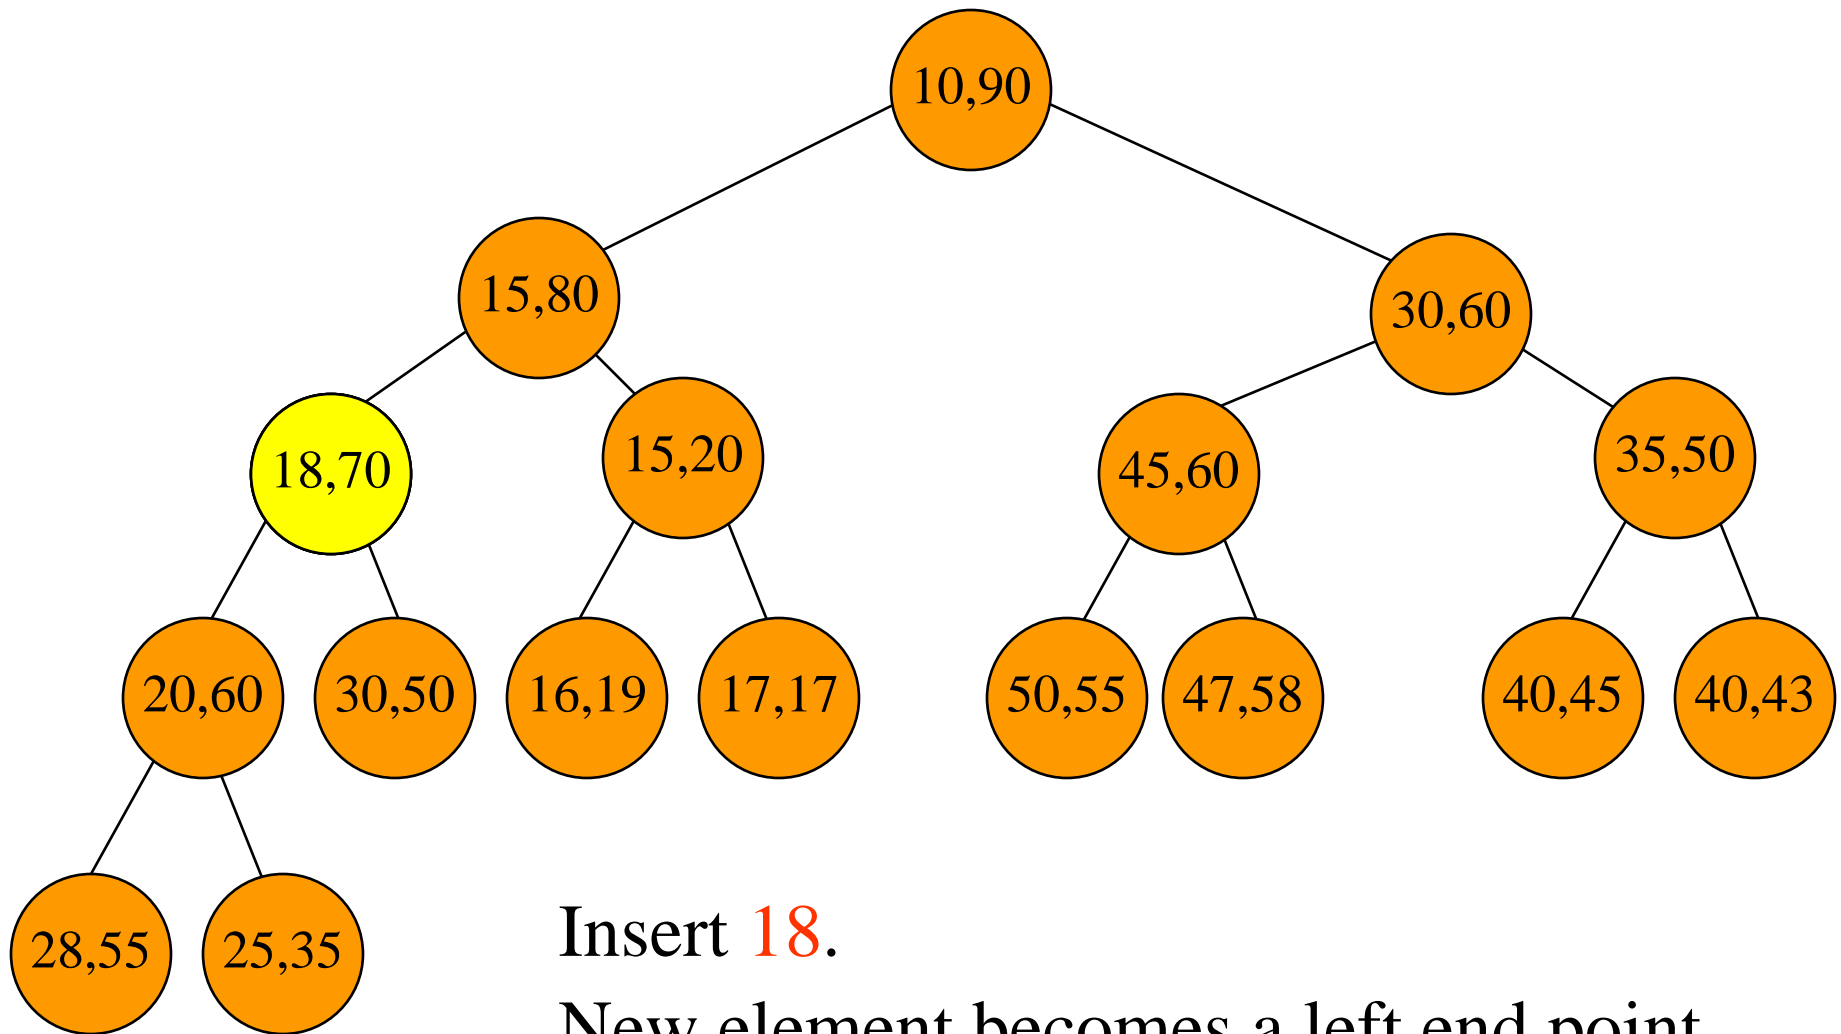


Insert 18.

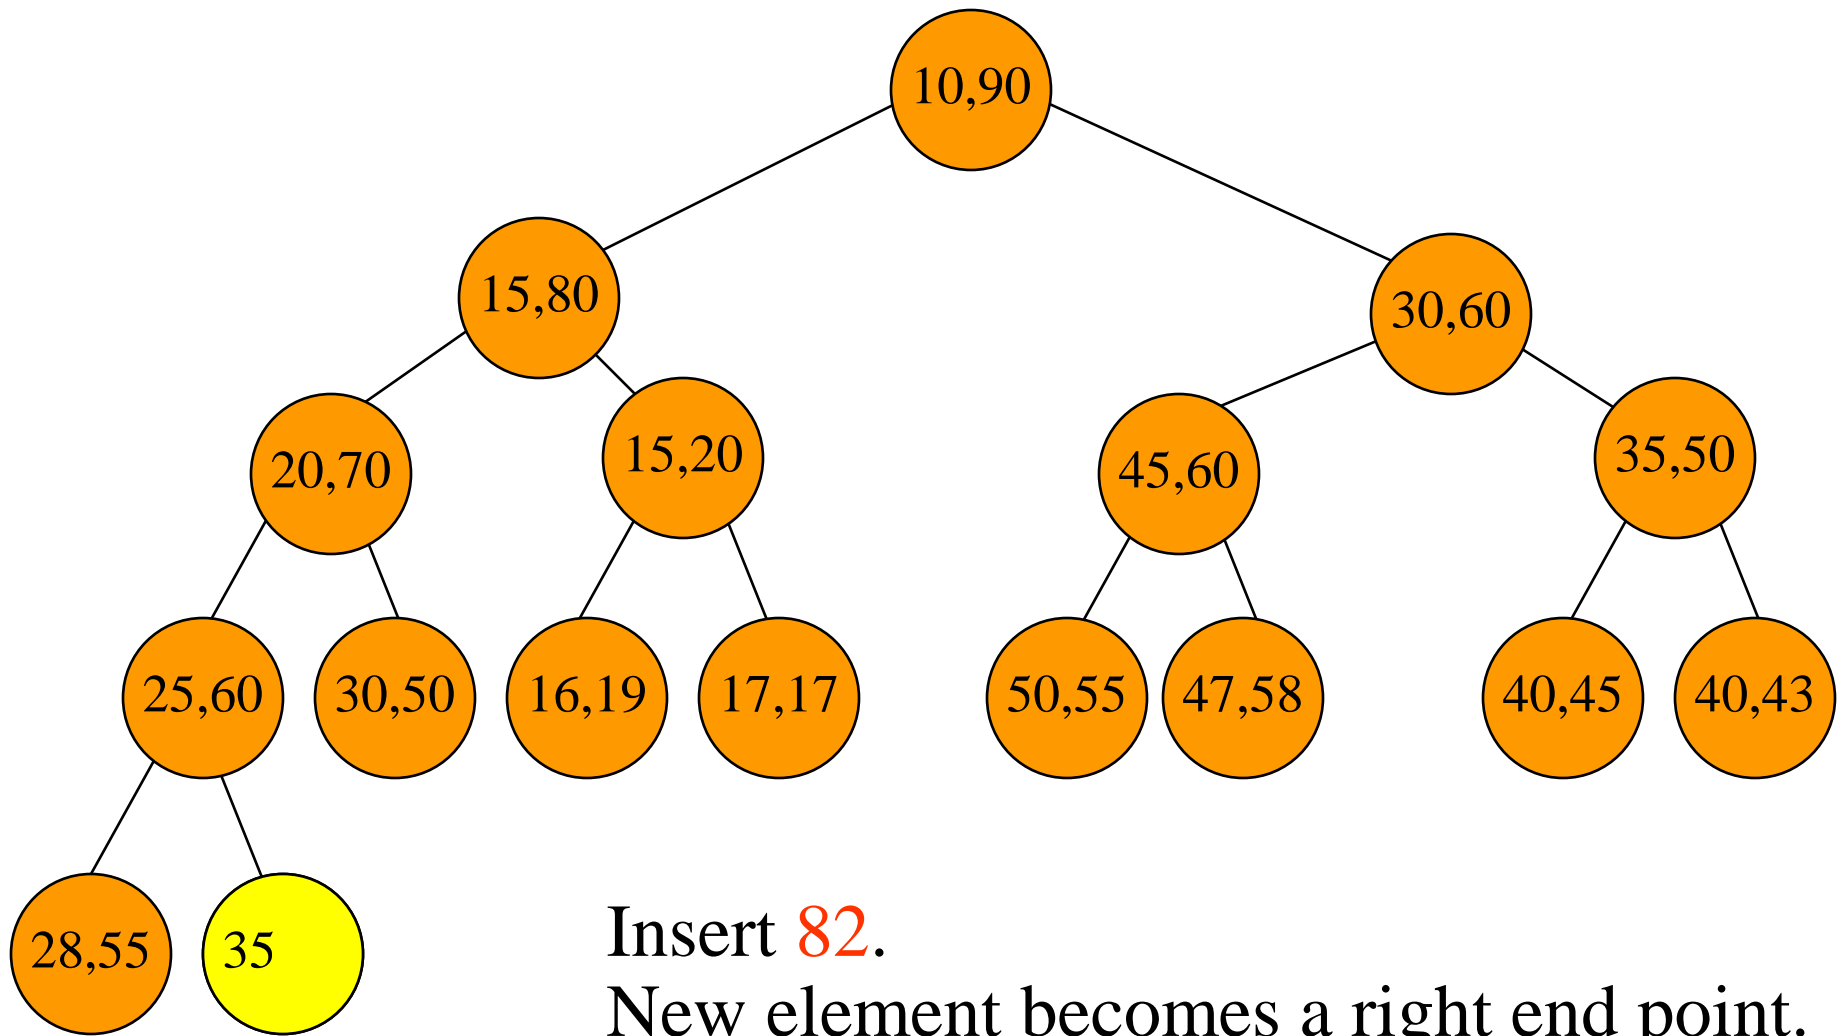
New element becomes a left end point.

Insert new element into min heap.

# Another Insert

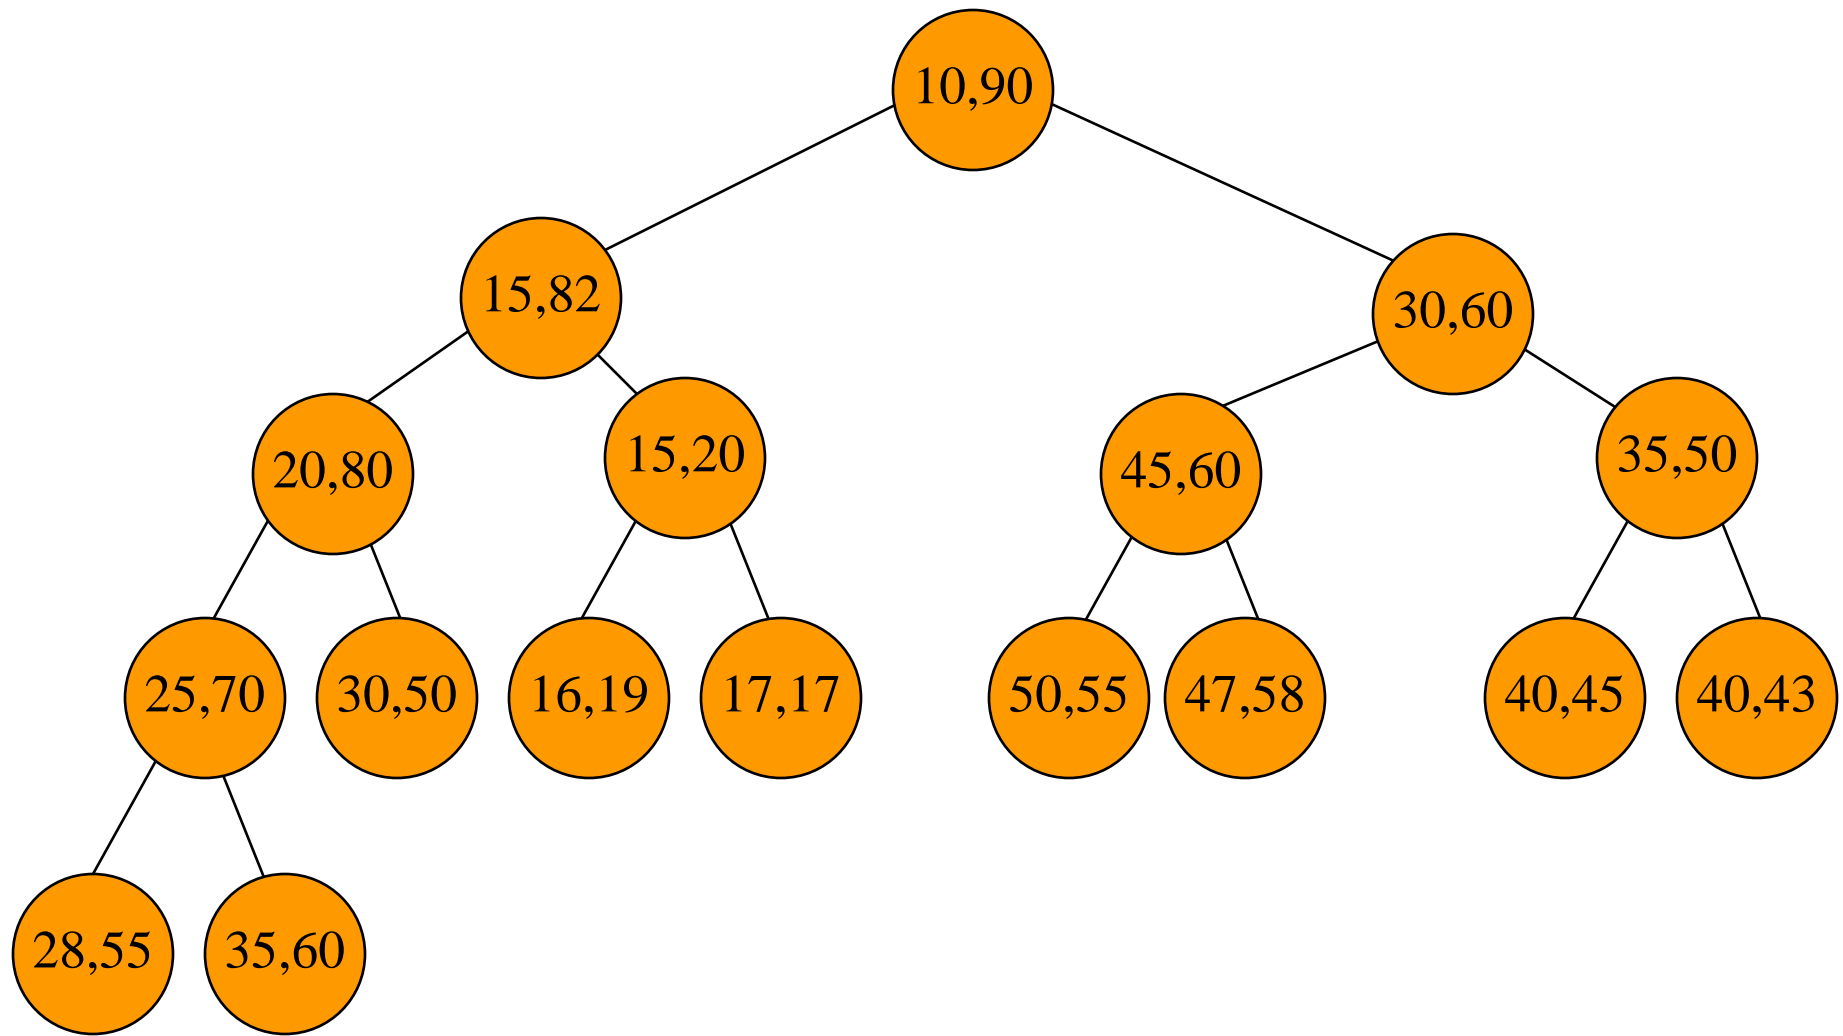


# Yet Another Insert






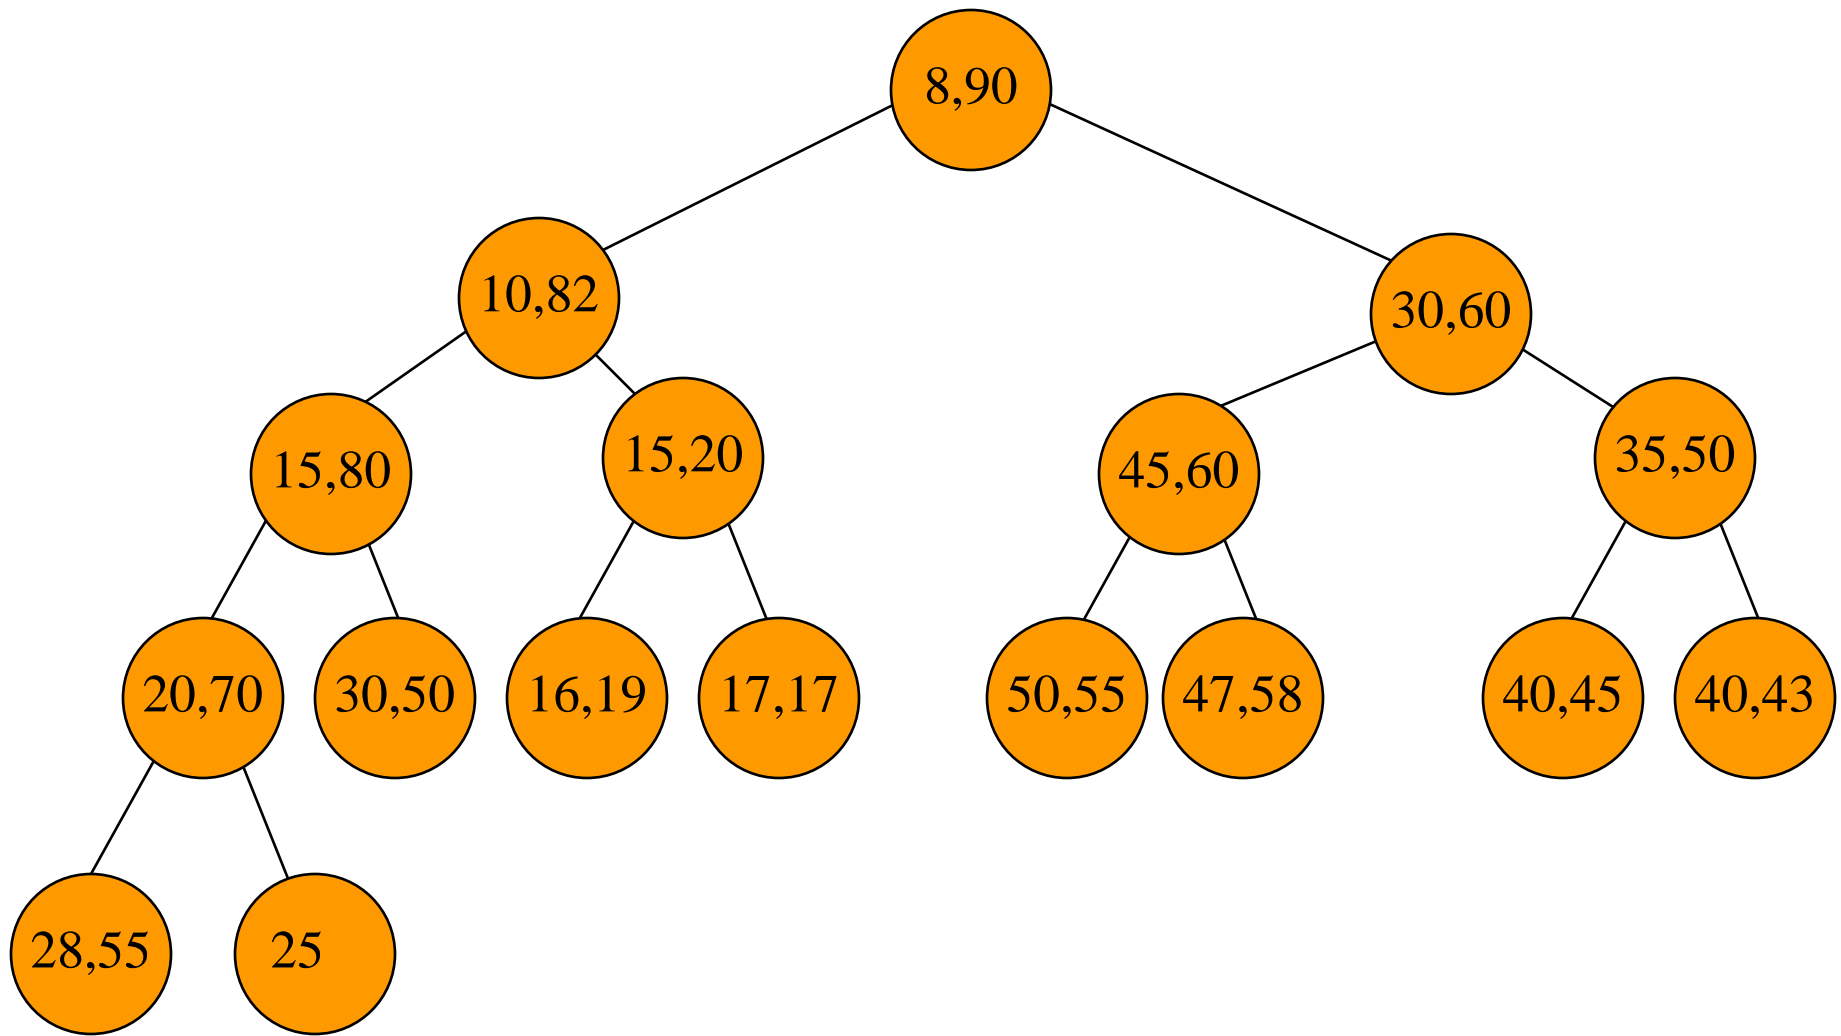
# After 82 Inserted



# One More Insert Example



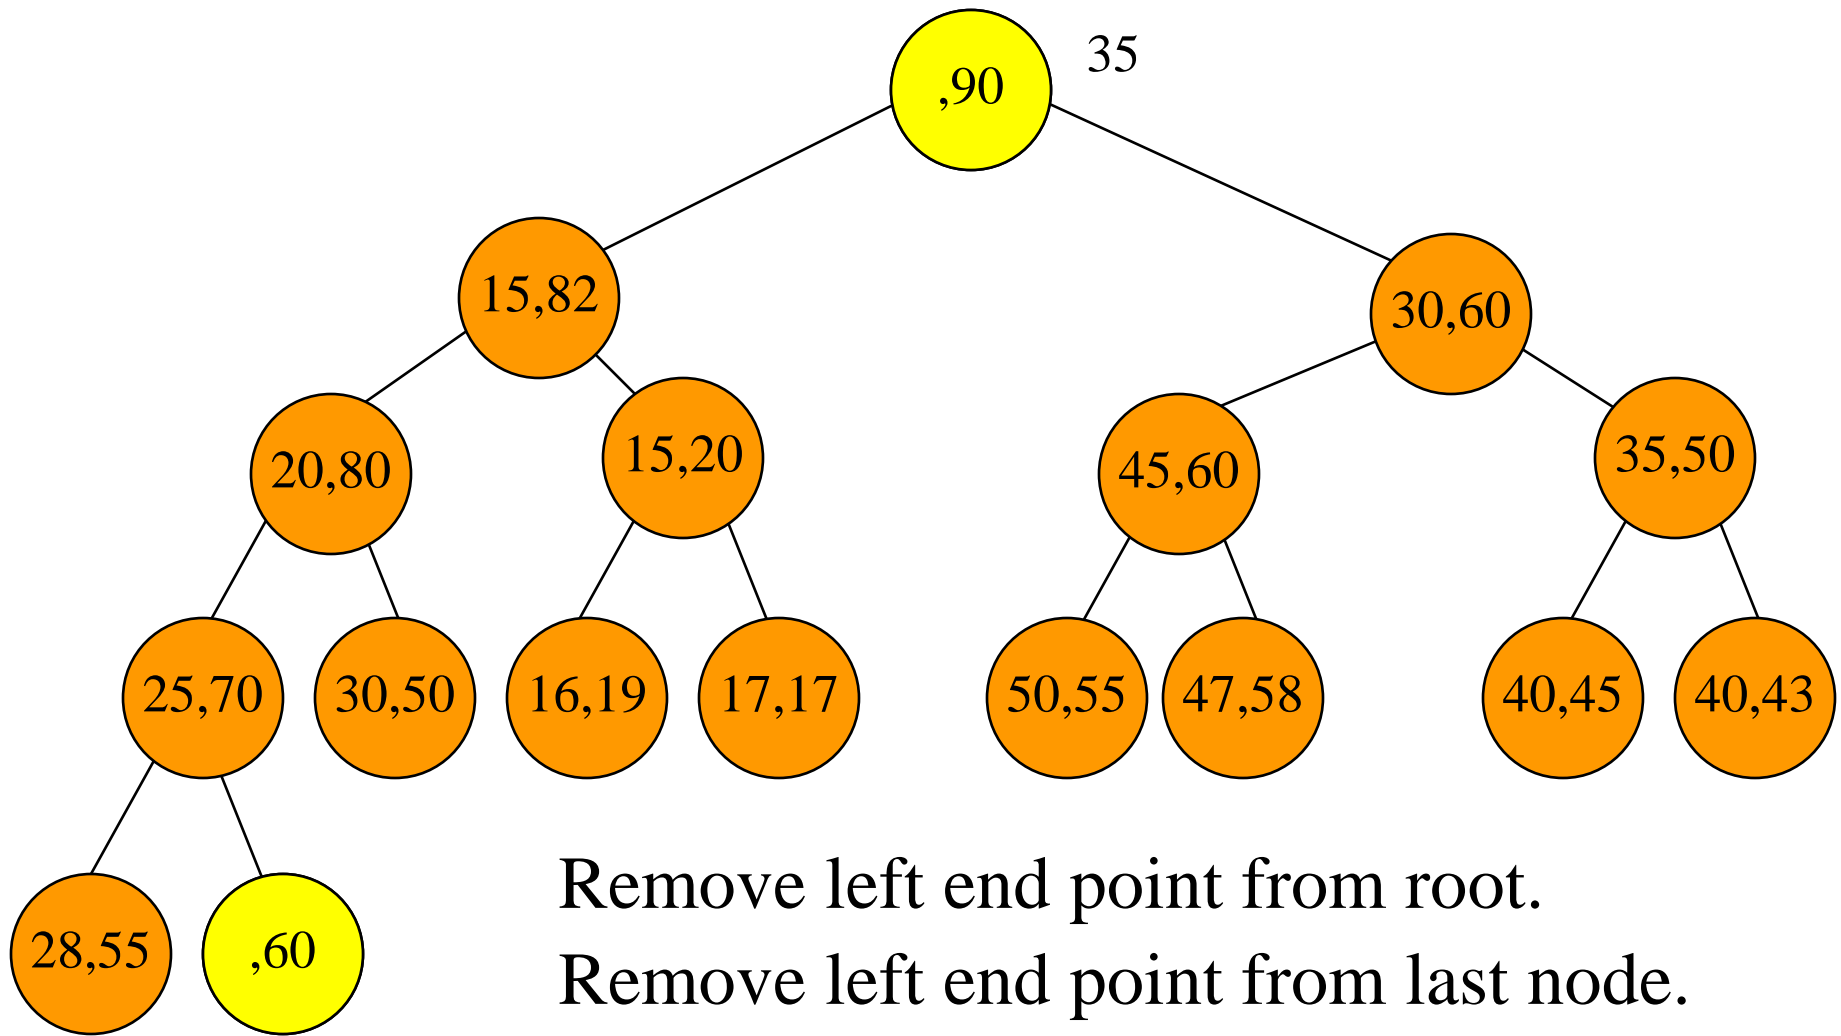
# After 8 Is Inserted



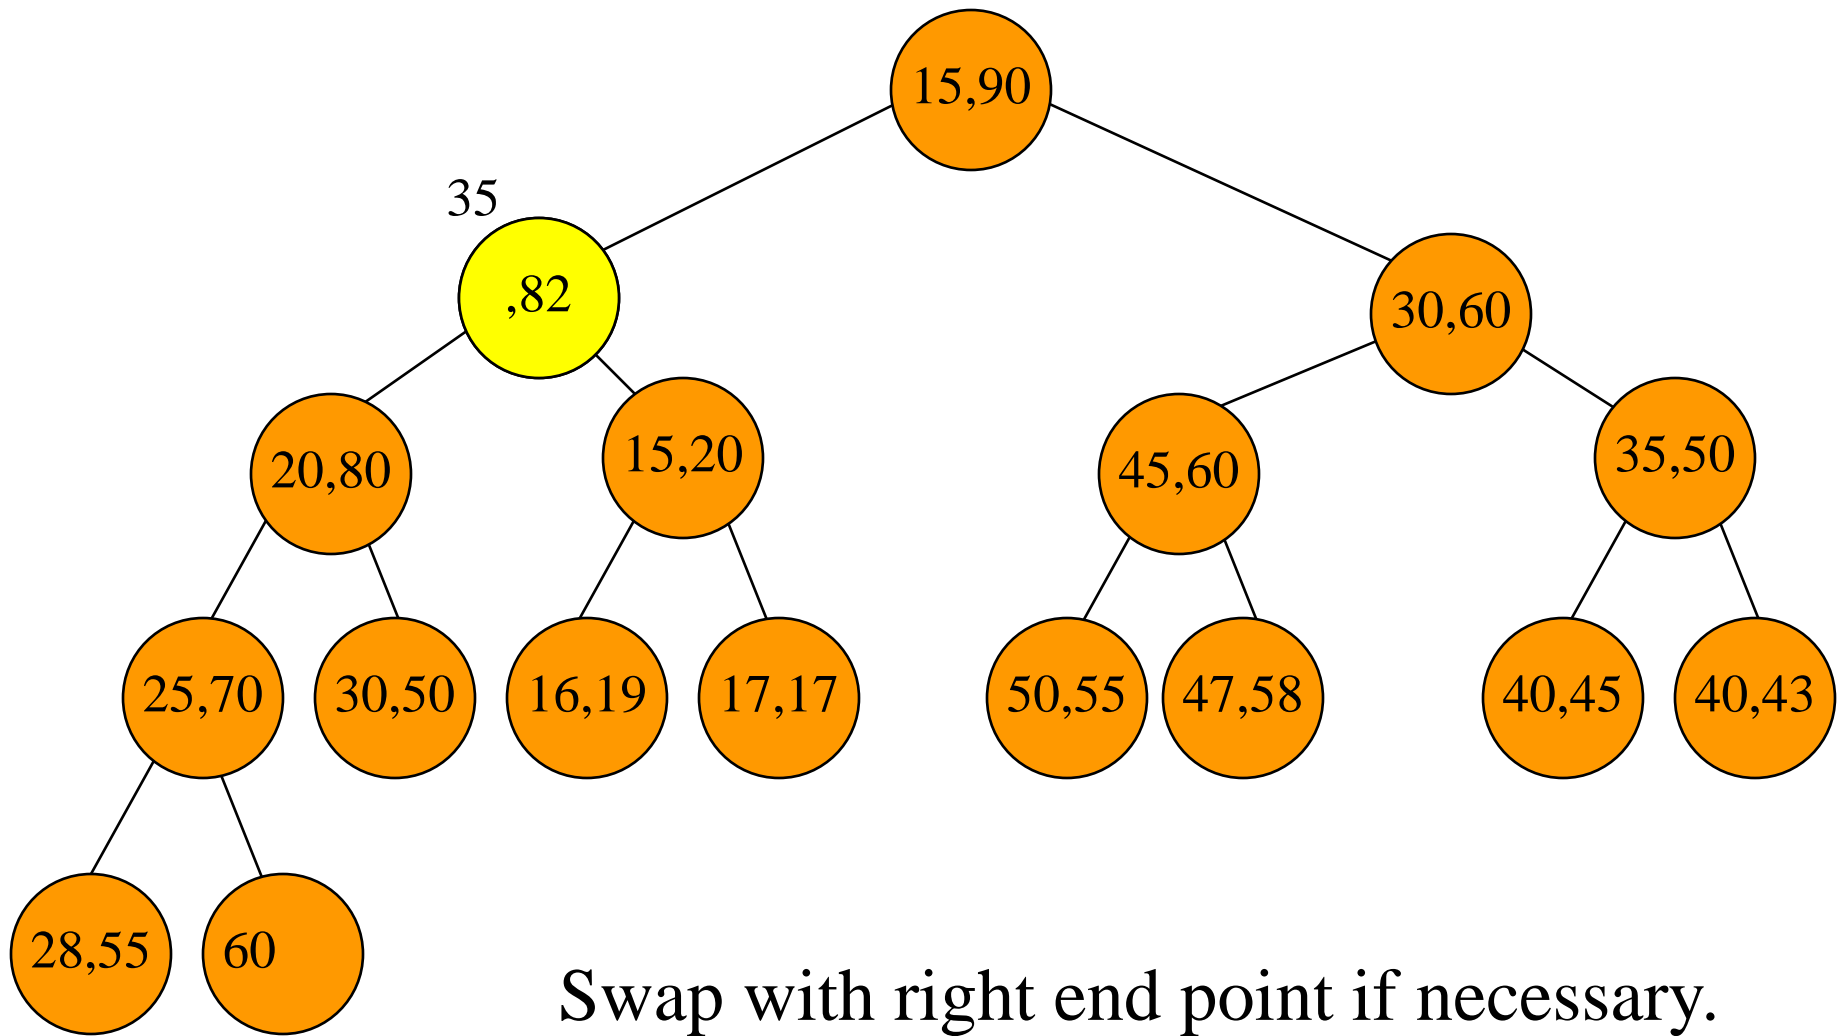
# Remove Min Element

- $n = 0 \Rightarrow$  fail.
- $n = 1 \Rightarrow$  heap becomes empty.
- $n = 2 \Rightarrow$  only one node, take out left end point.
- $n > 2 \Rightarrow$  not as simple.

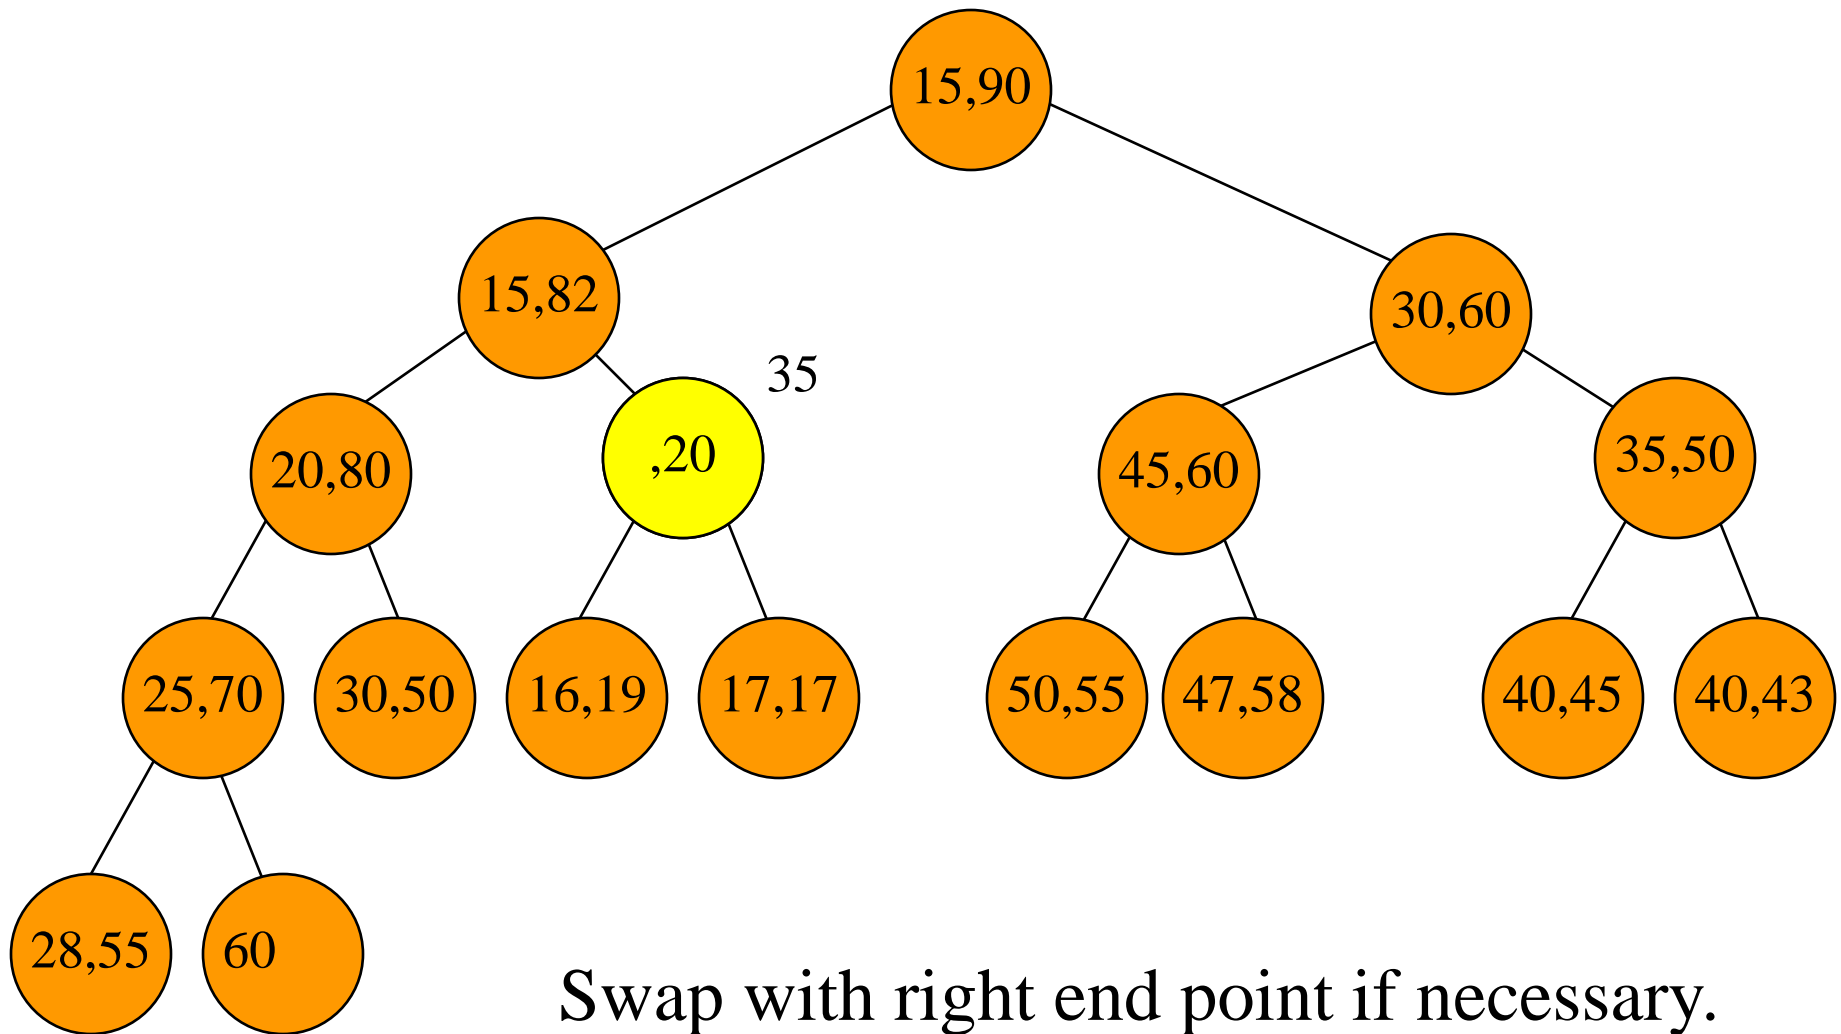
# Remove Min Element Example



# Remove Min Element Example

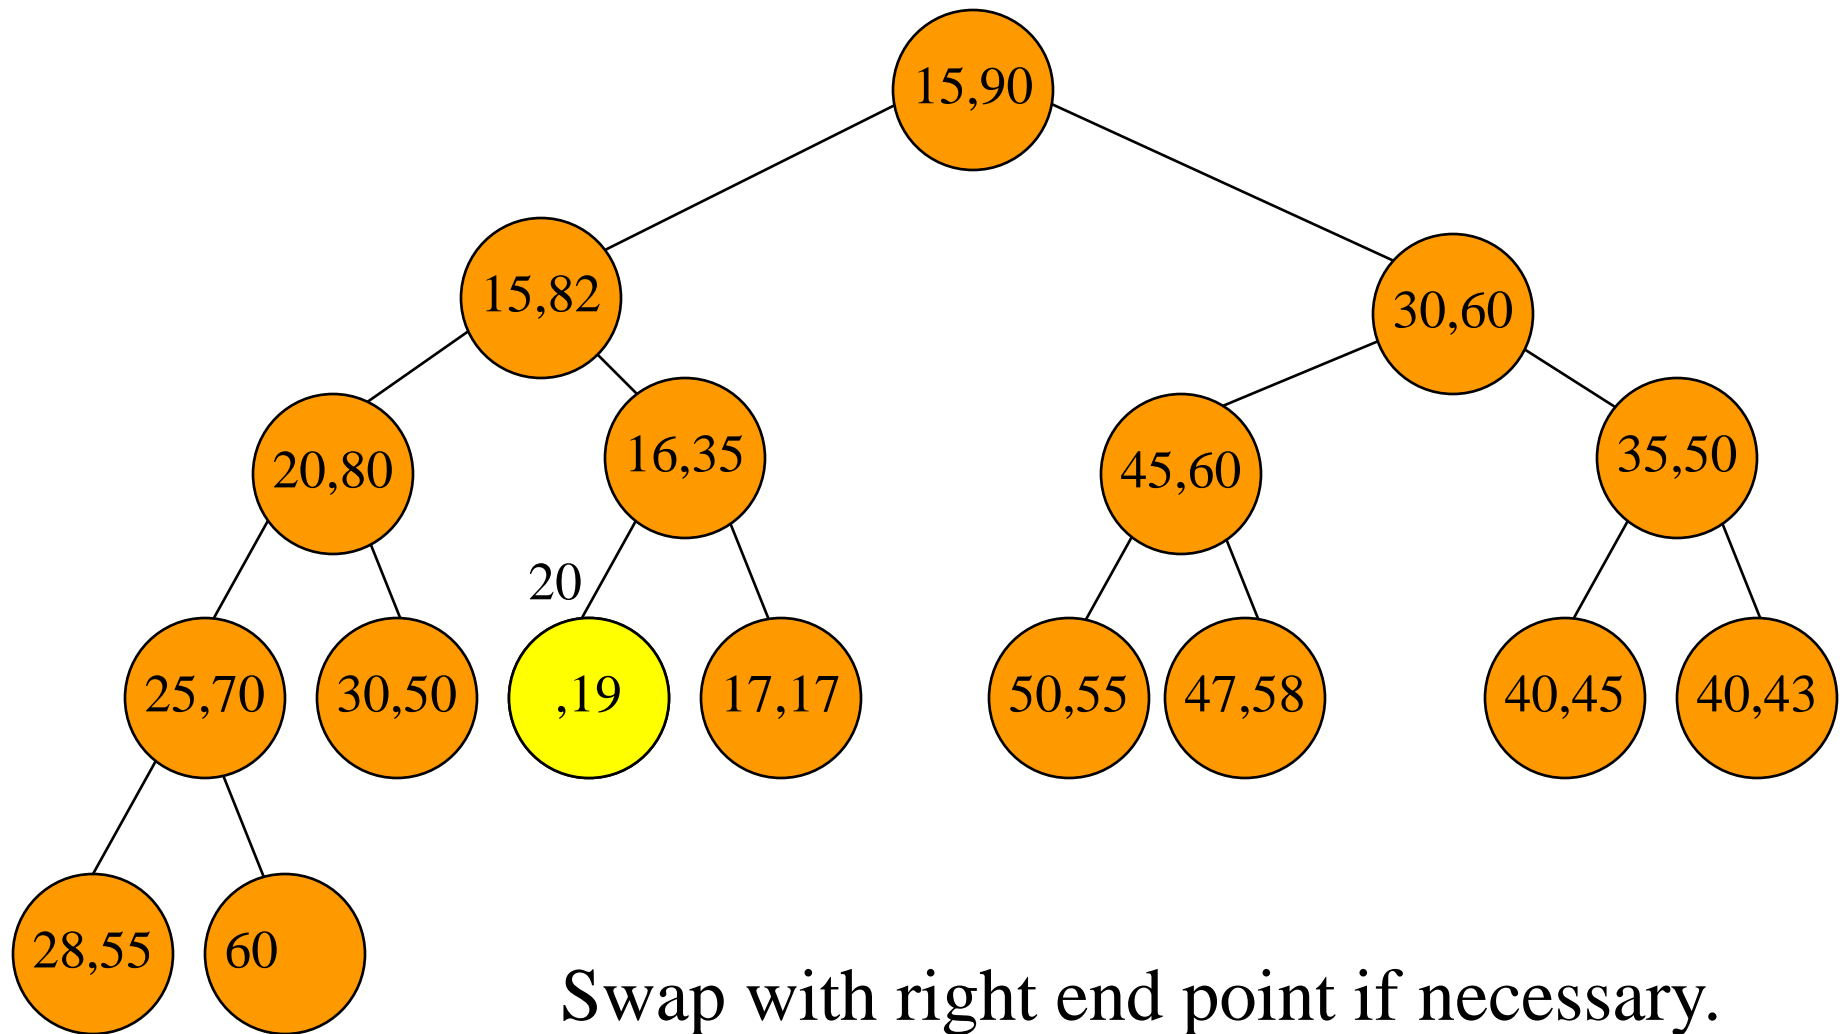


# Remove Min Element Example



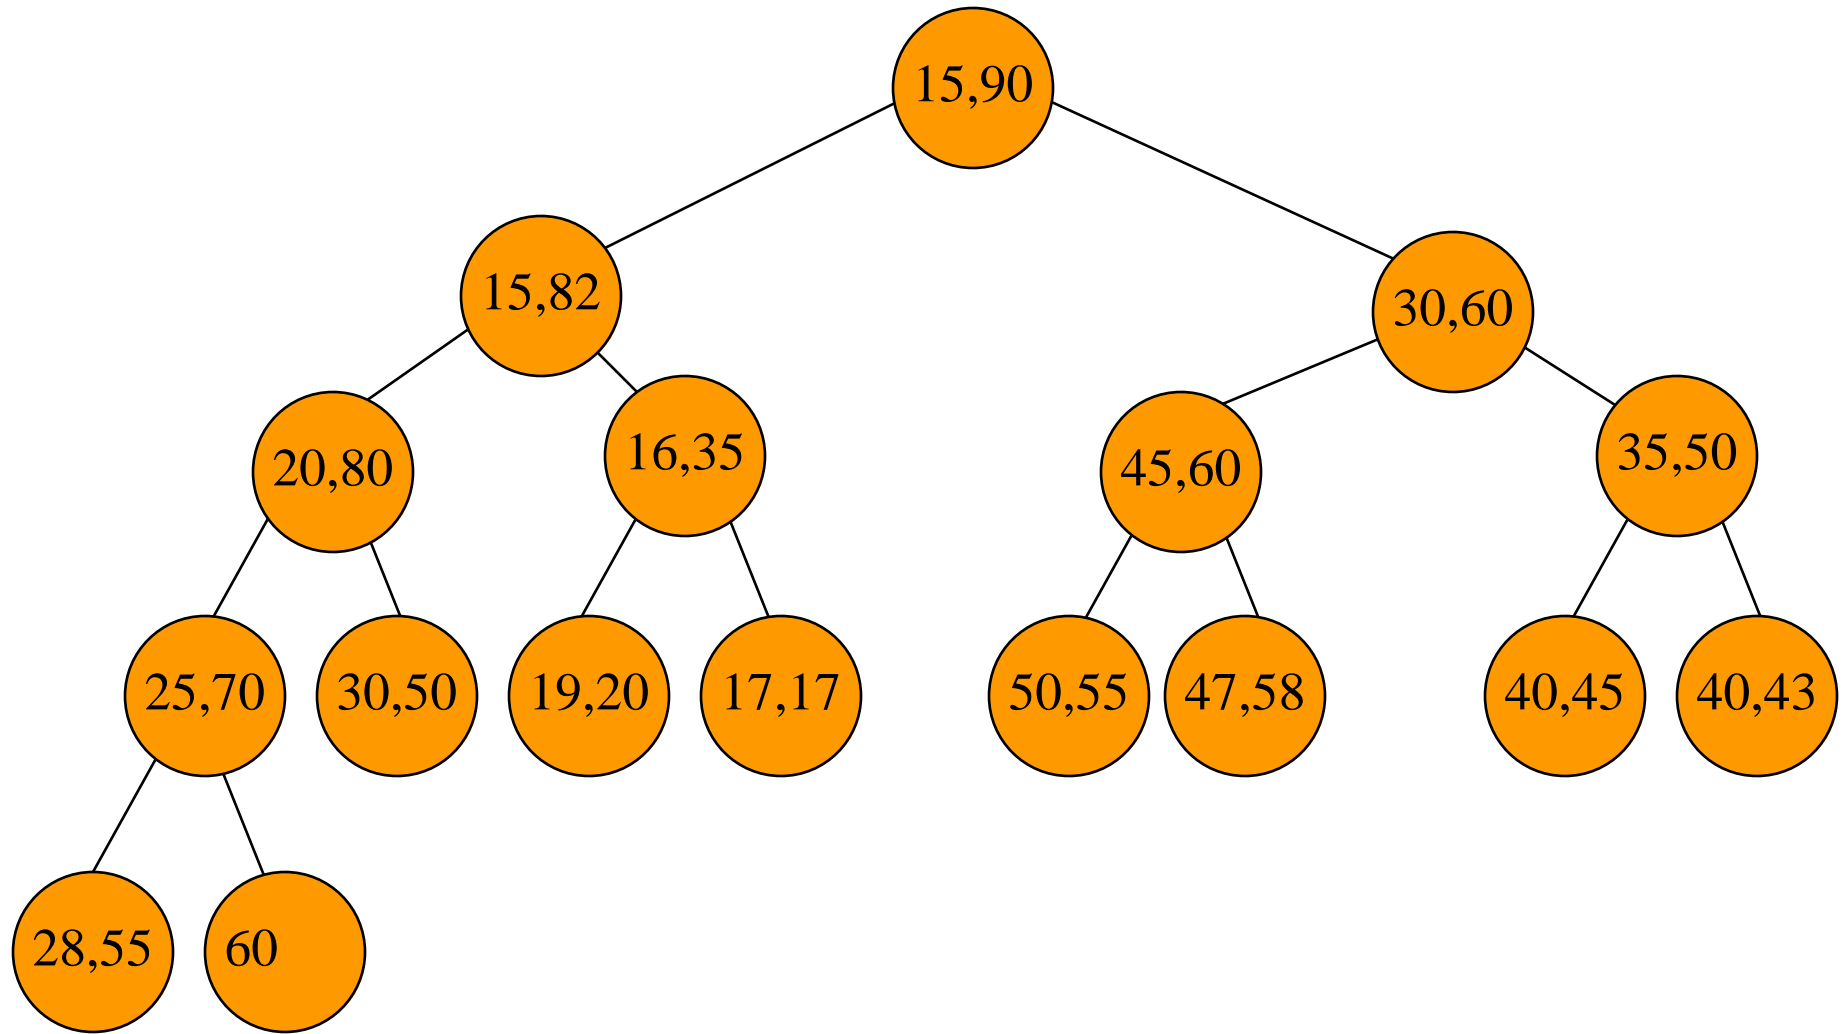
Swap with right end point if necessary.

# Remove Min Element Example

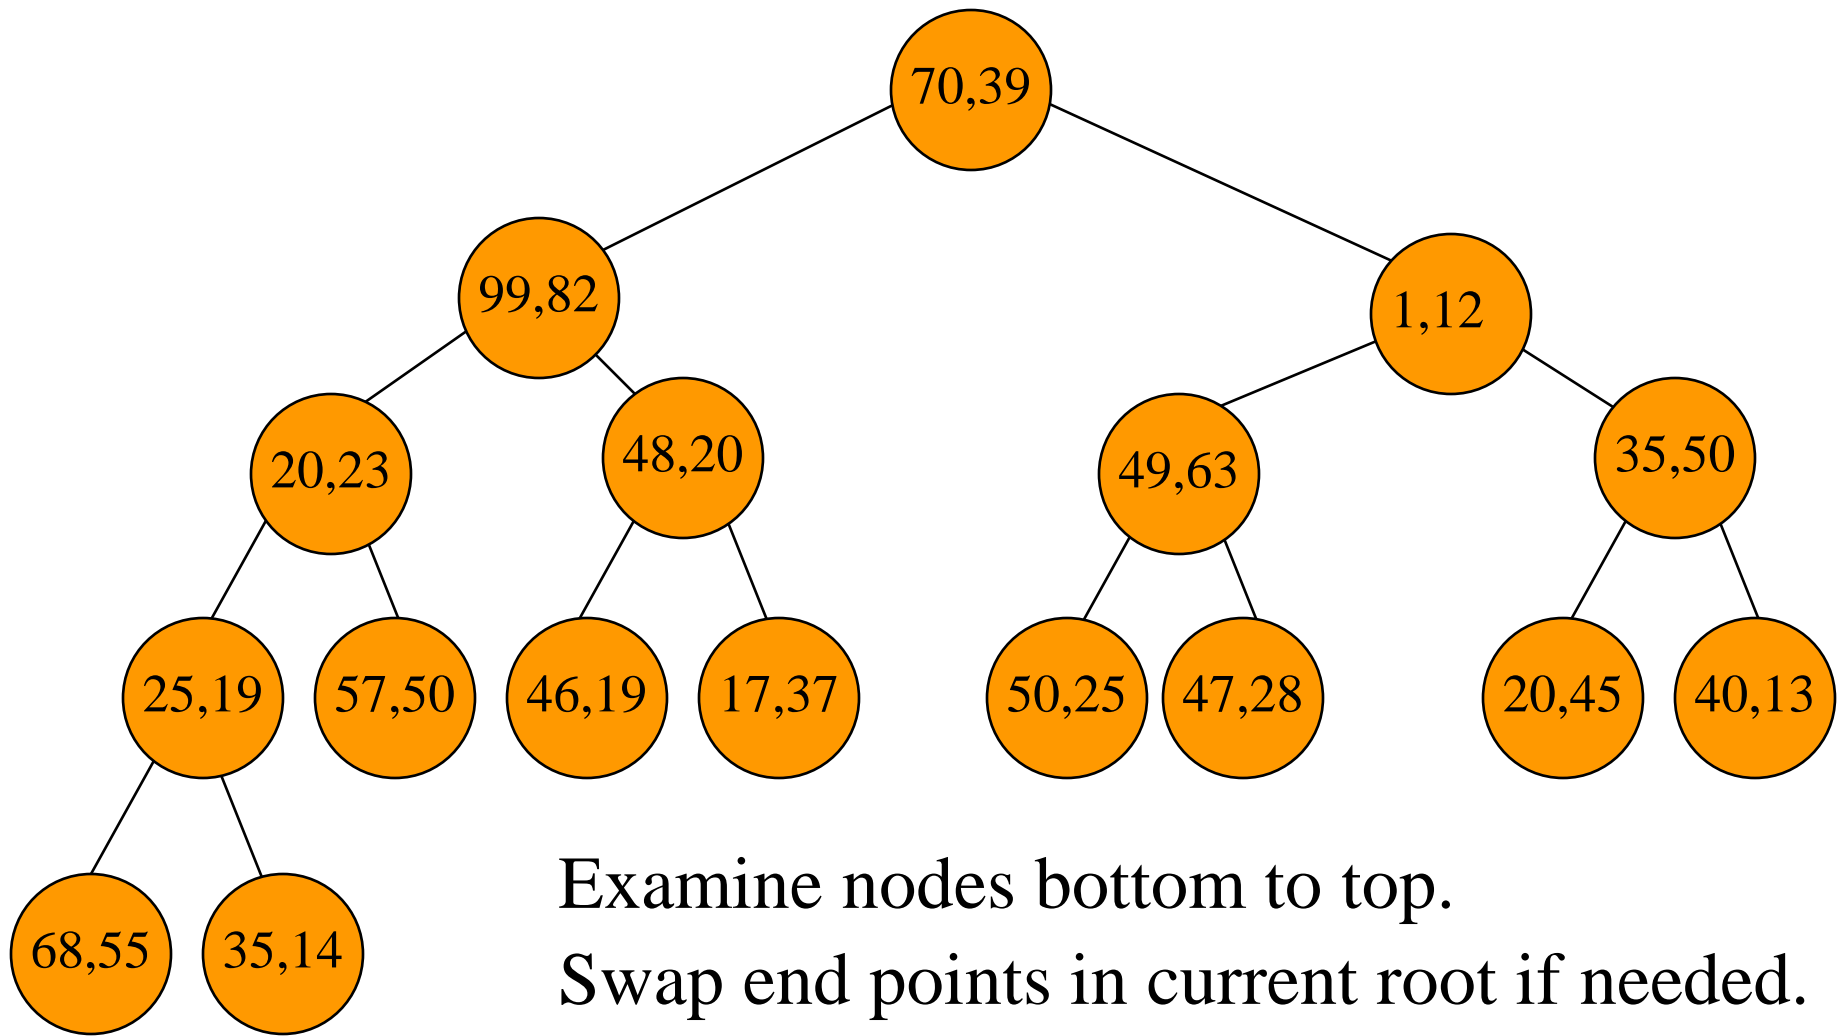




# Remove Min Element Example



# Initialize



# Hints

- You can use Min-max heap or Deap to implement Double-ended Priority Queue
- Count disk I/Os (File reads/writes)

# Phase 3

External Sort: Merge Sort


# Internal Merge Sort Review

- Phase 1
  - Create initial sorted segments
    - Natural segments
    - Insertion sort
- Phase 2
  - Merge pairs of sorted segments, in merge passes, until only **1** segment remains.

# External Merge Sort

- Two phases.
  - Run generation.
    - A run is a sorted sequence of records.
  - Run merging.

# Run Generation




- Input blocks.
- Sort.
- Output as a run.

# Run Merging


- Merge Pass.
  - Pairwise merge the (initial) runs.
  - In a **merge pass** all runs (except possibly one) are pairwise merged.
- Perform multiple merge passes, reducing the number of runs to **1**.



# Merge 20 Runs



# Merge R1 and R2




- Fill **I0** (Input 0) from **R1** and **I1** from **R2**.
- Merge from **I0** and **I1** to output buffer.
- Write whenever output buffer full.
- Read whenever input buffer empty.

# Phase 4

Merge Sort: Improve Run Generation

# Improve Run Generation

- Overlap input, output, and internal sorting.




# Improve Run Generation

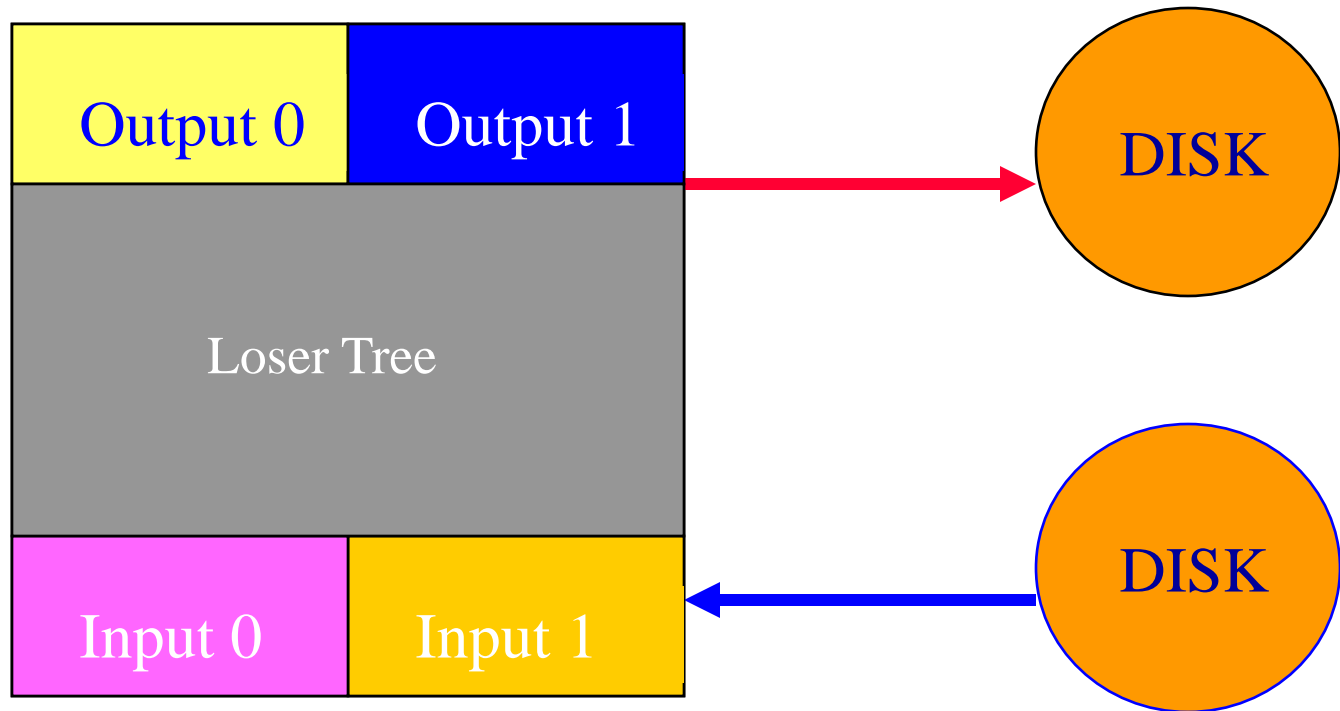
- Generate runs whose length (on average) exceeds memory size.
- Equivalent to reducing number of runs generated.

# Improve Run Generation

- Overlap input,output, and internal CPU work.
- Reduce the number of runs (equivalently, increase average run length).

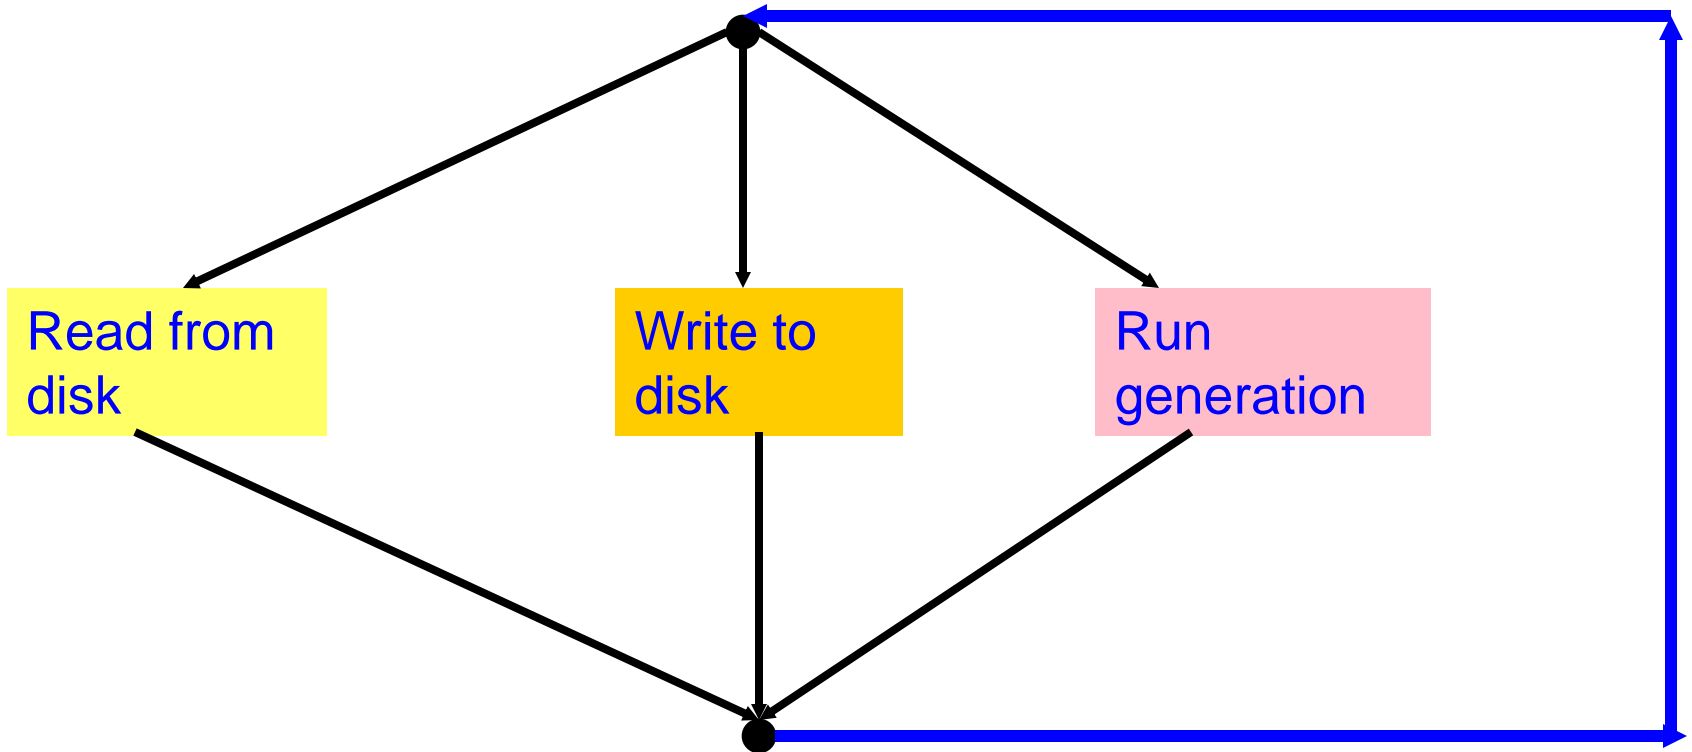


# New Strategy



- Use **2** input and **2** output buffers.
- Rest of memory is used for a min loser tree.
- Actually, 3 buffers adequate.


# Steady State Operation




- Synchronization is done when the active input buffer gets empty (the active output buffer will be full at this time).



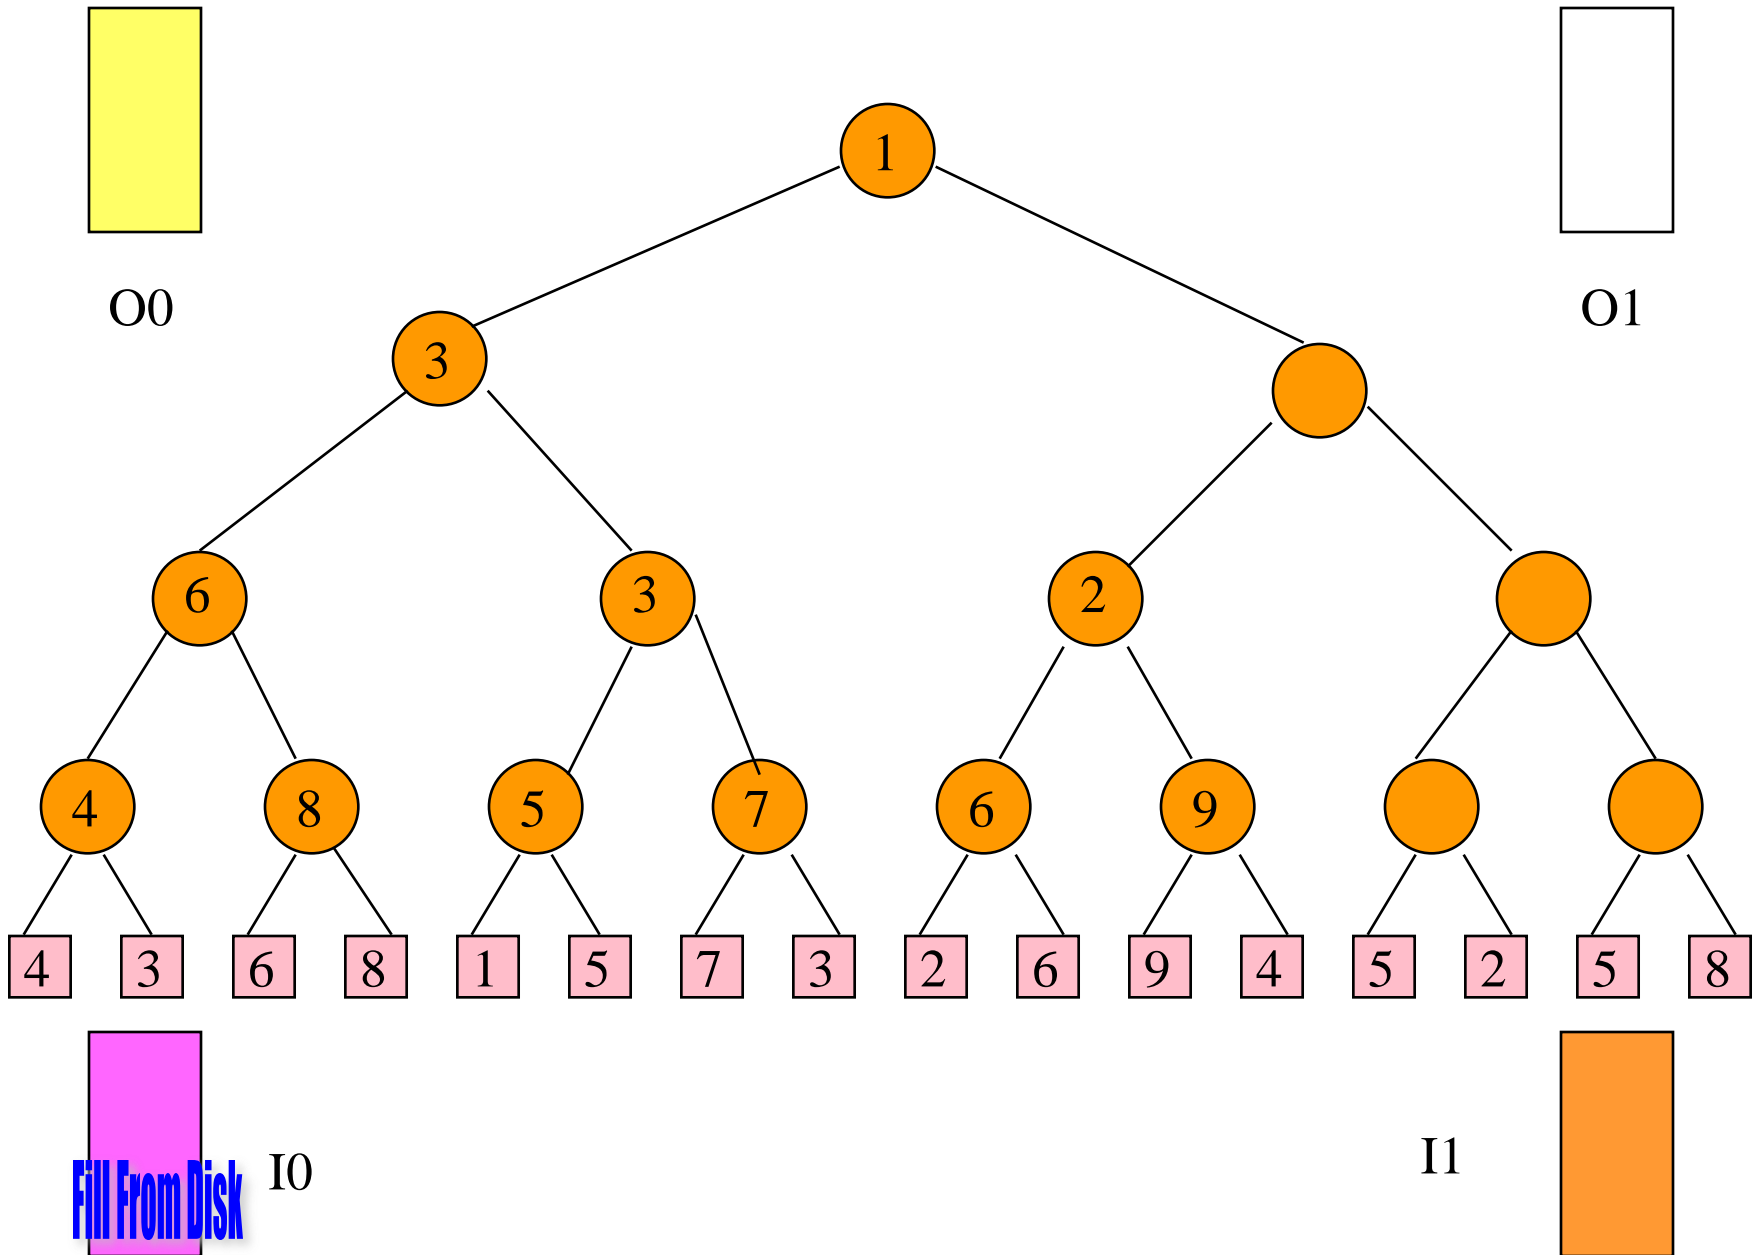
# Initialize



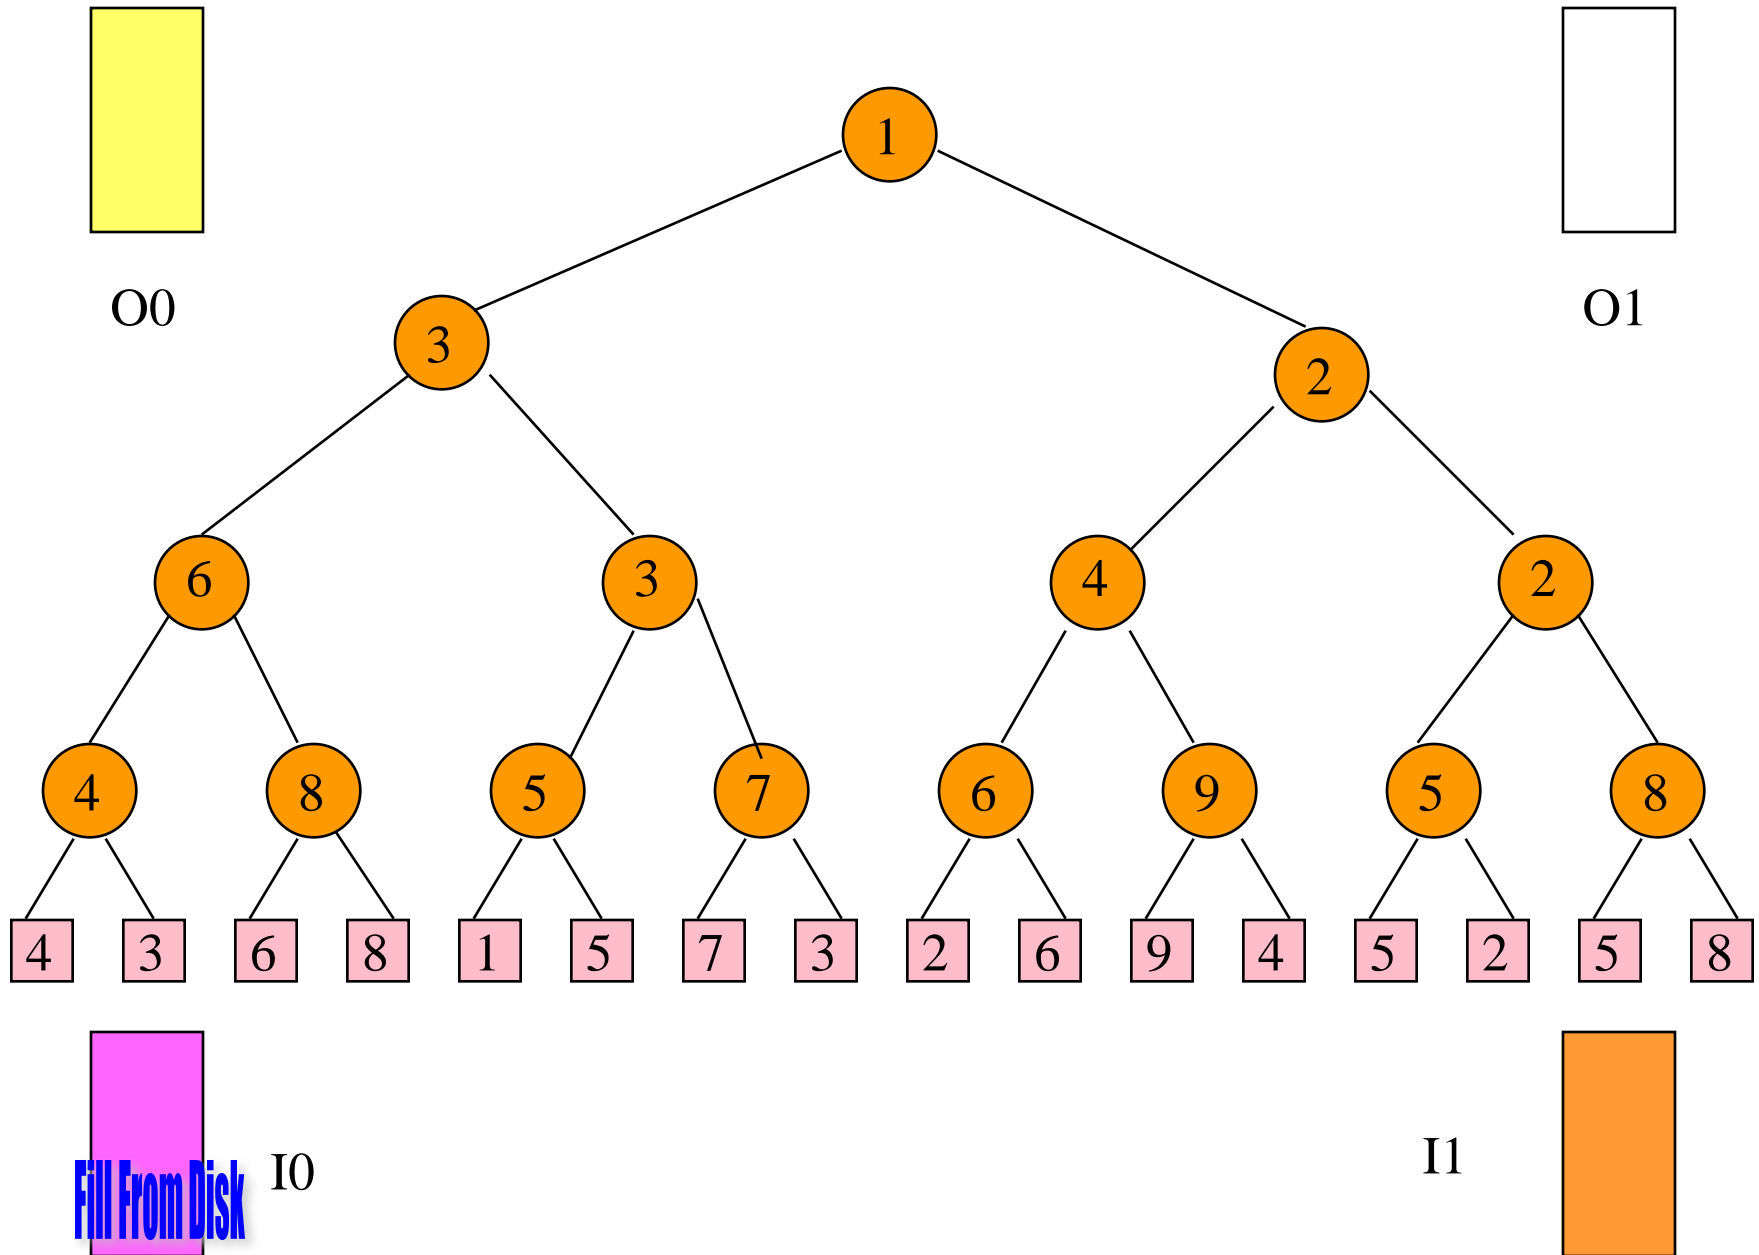
# Initialize



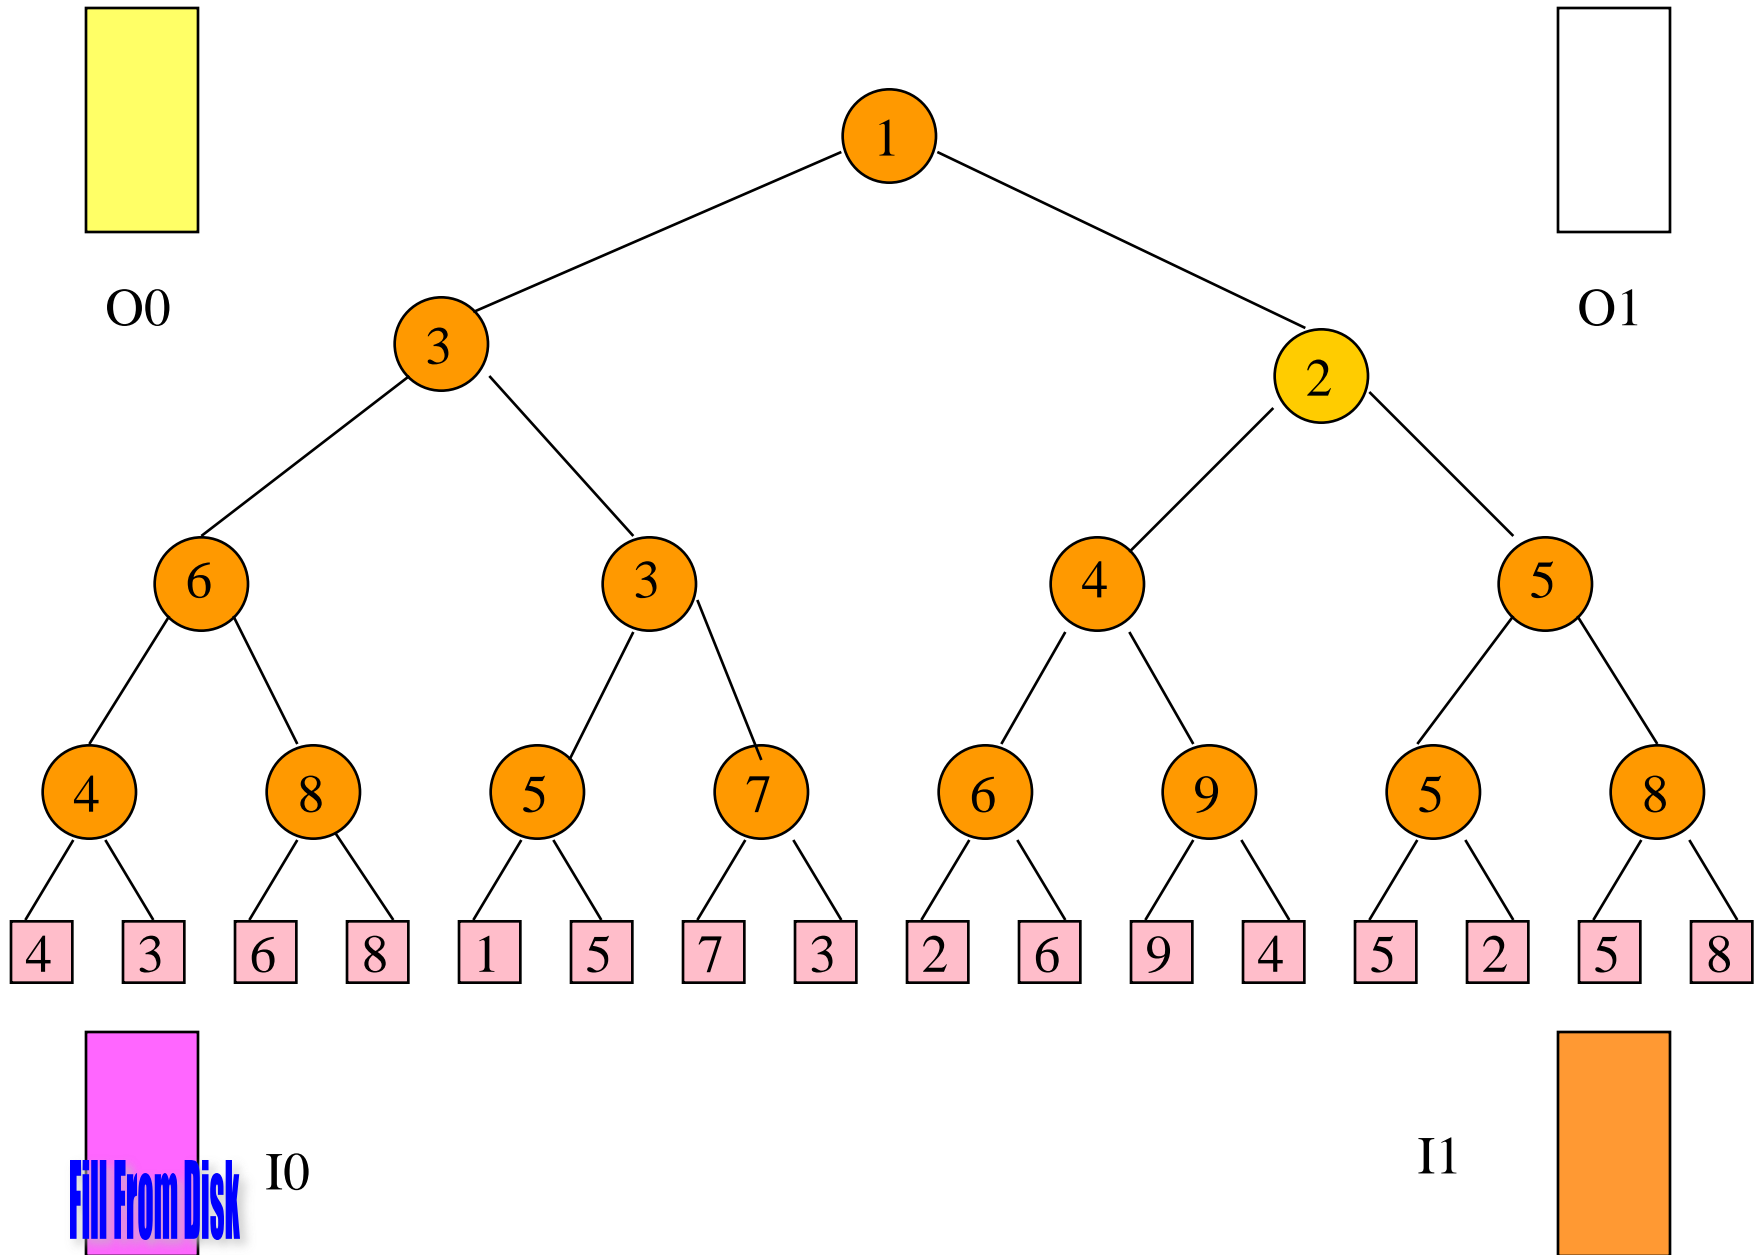
# Initialize



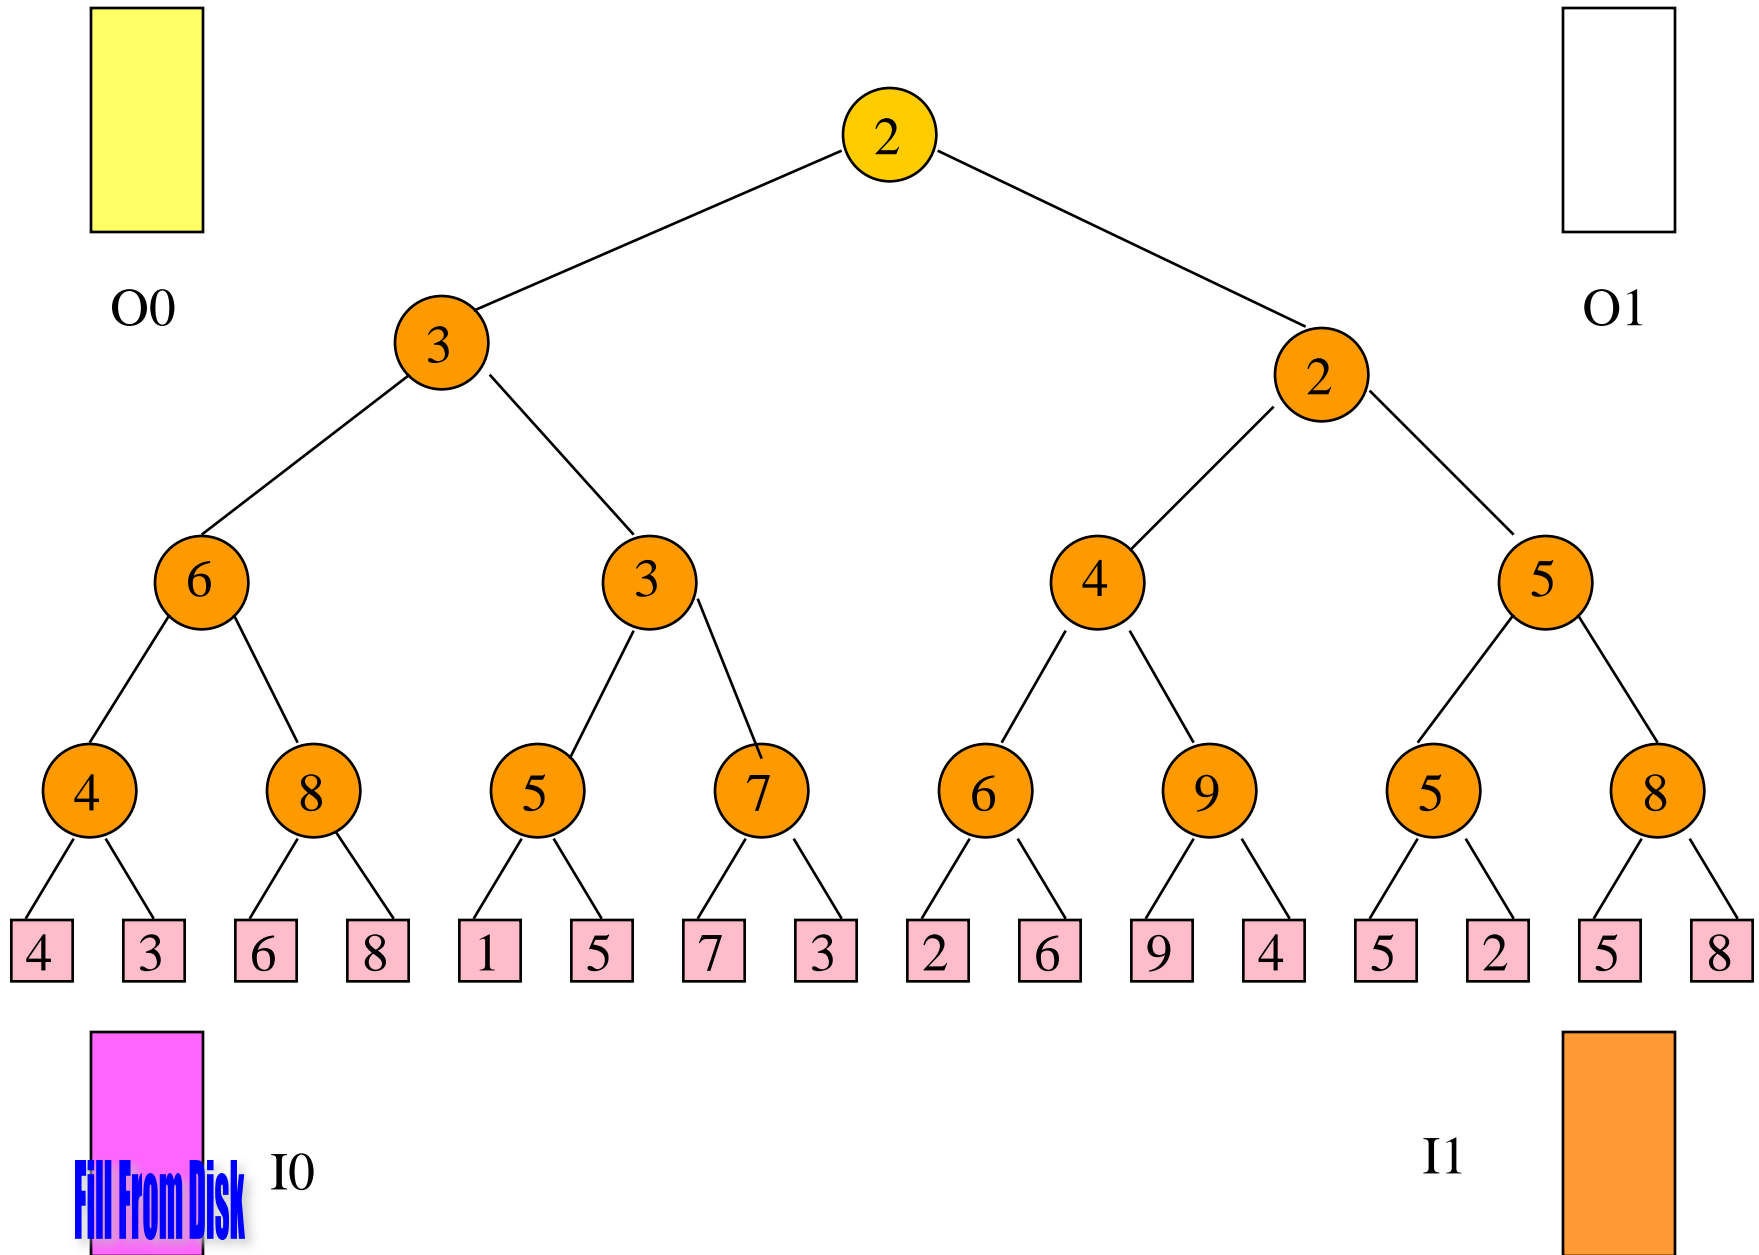
# Initialize



# Initialize

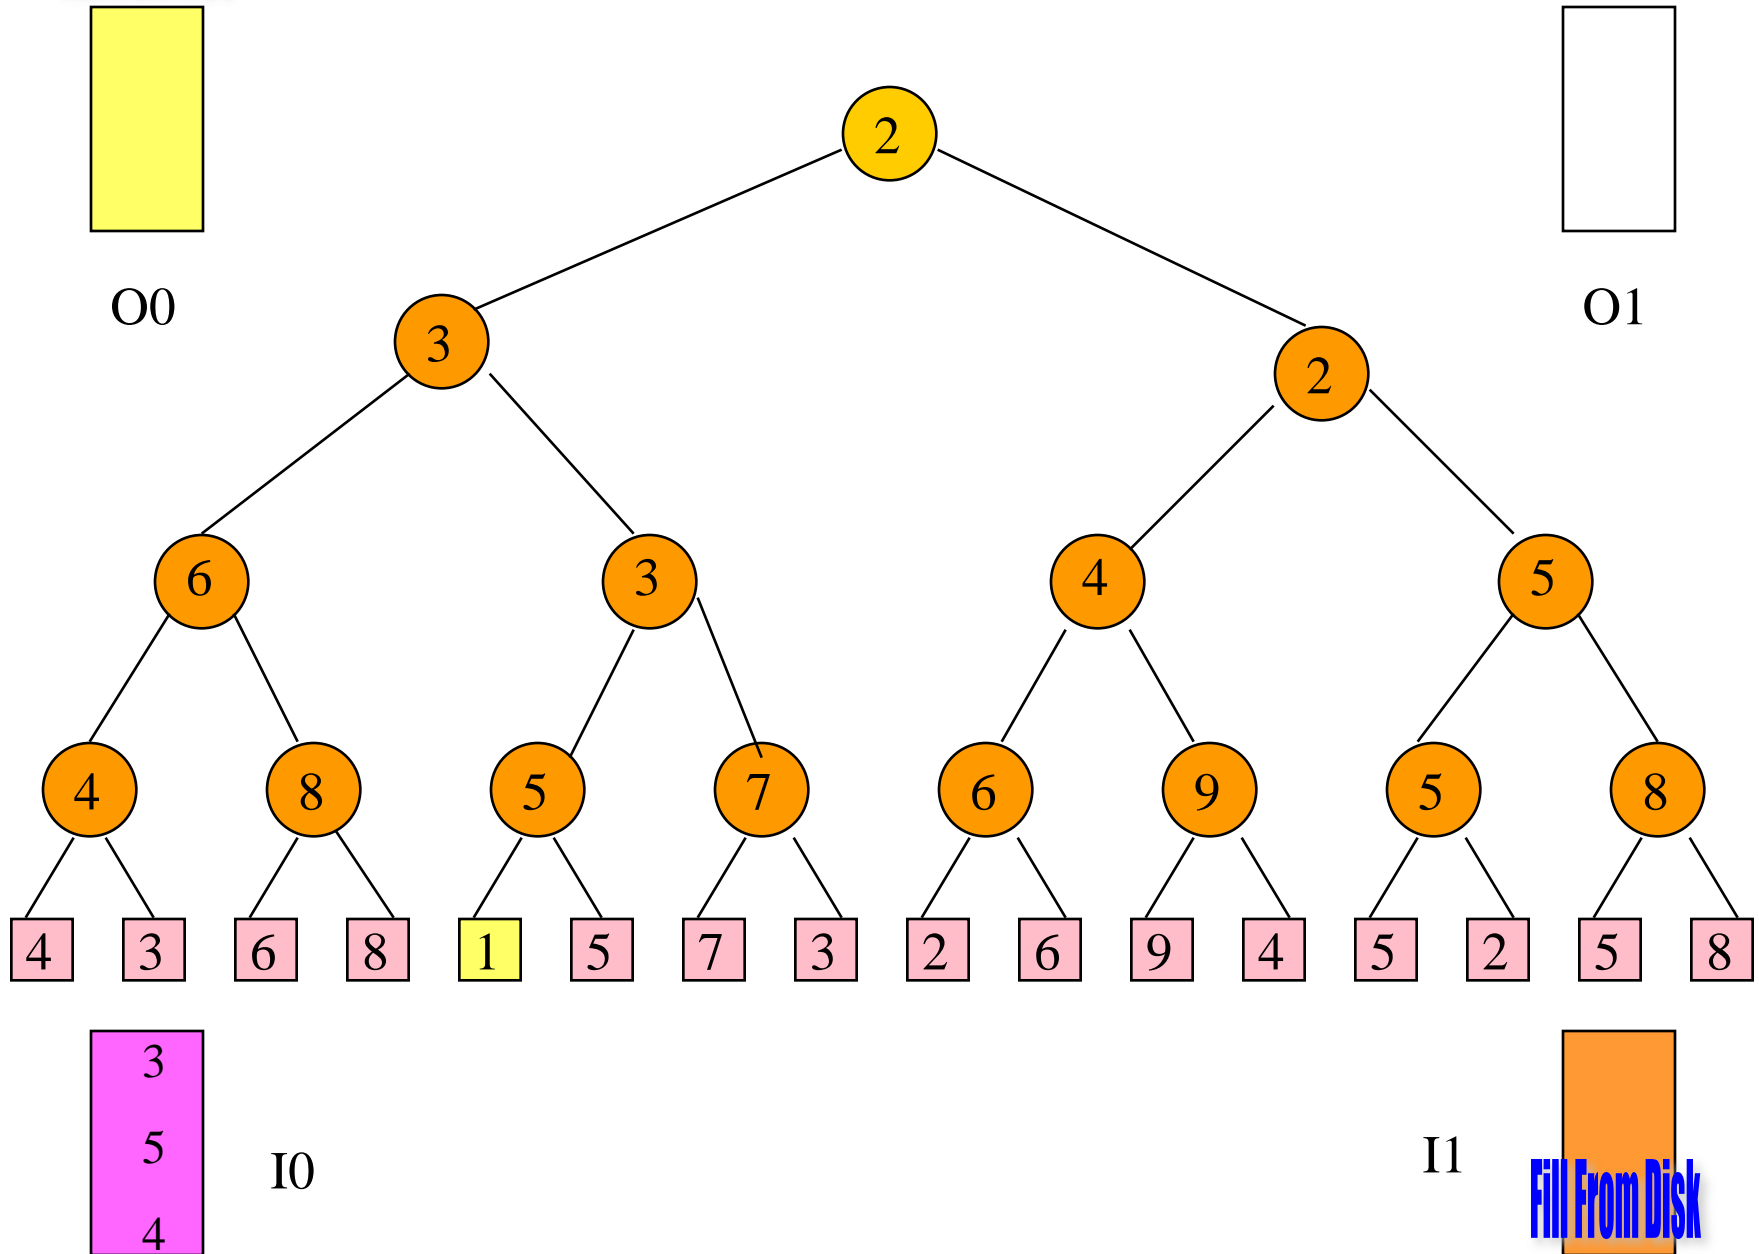


# Initialize



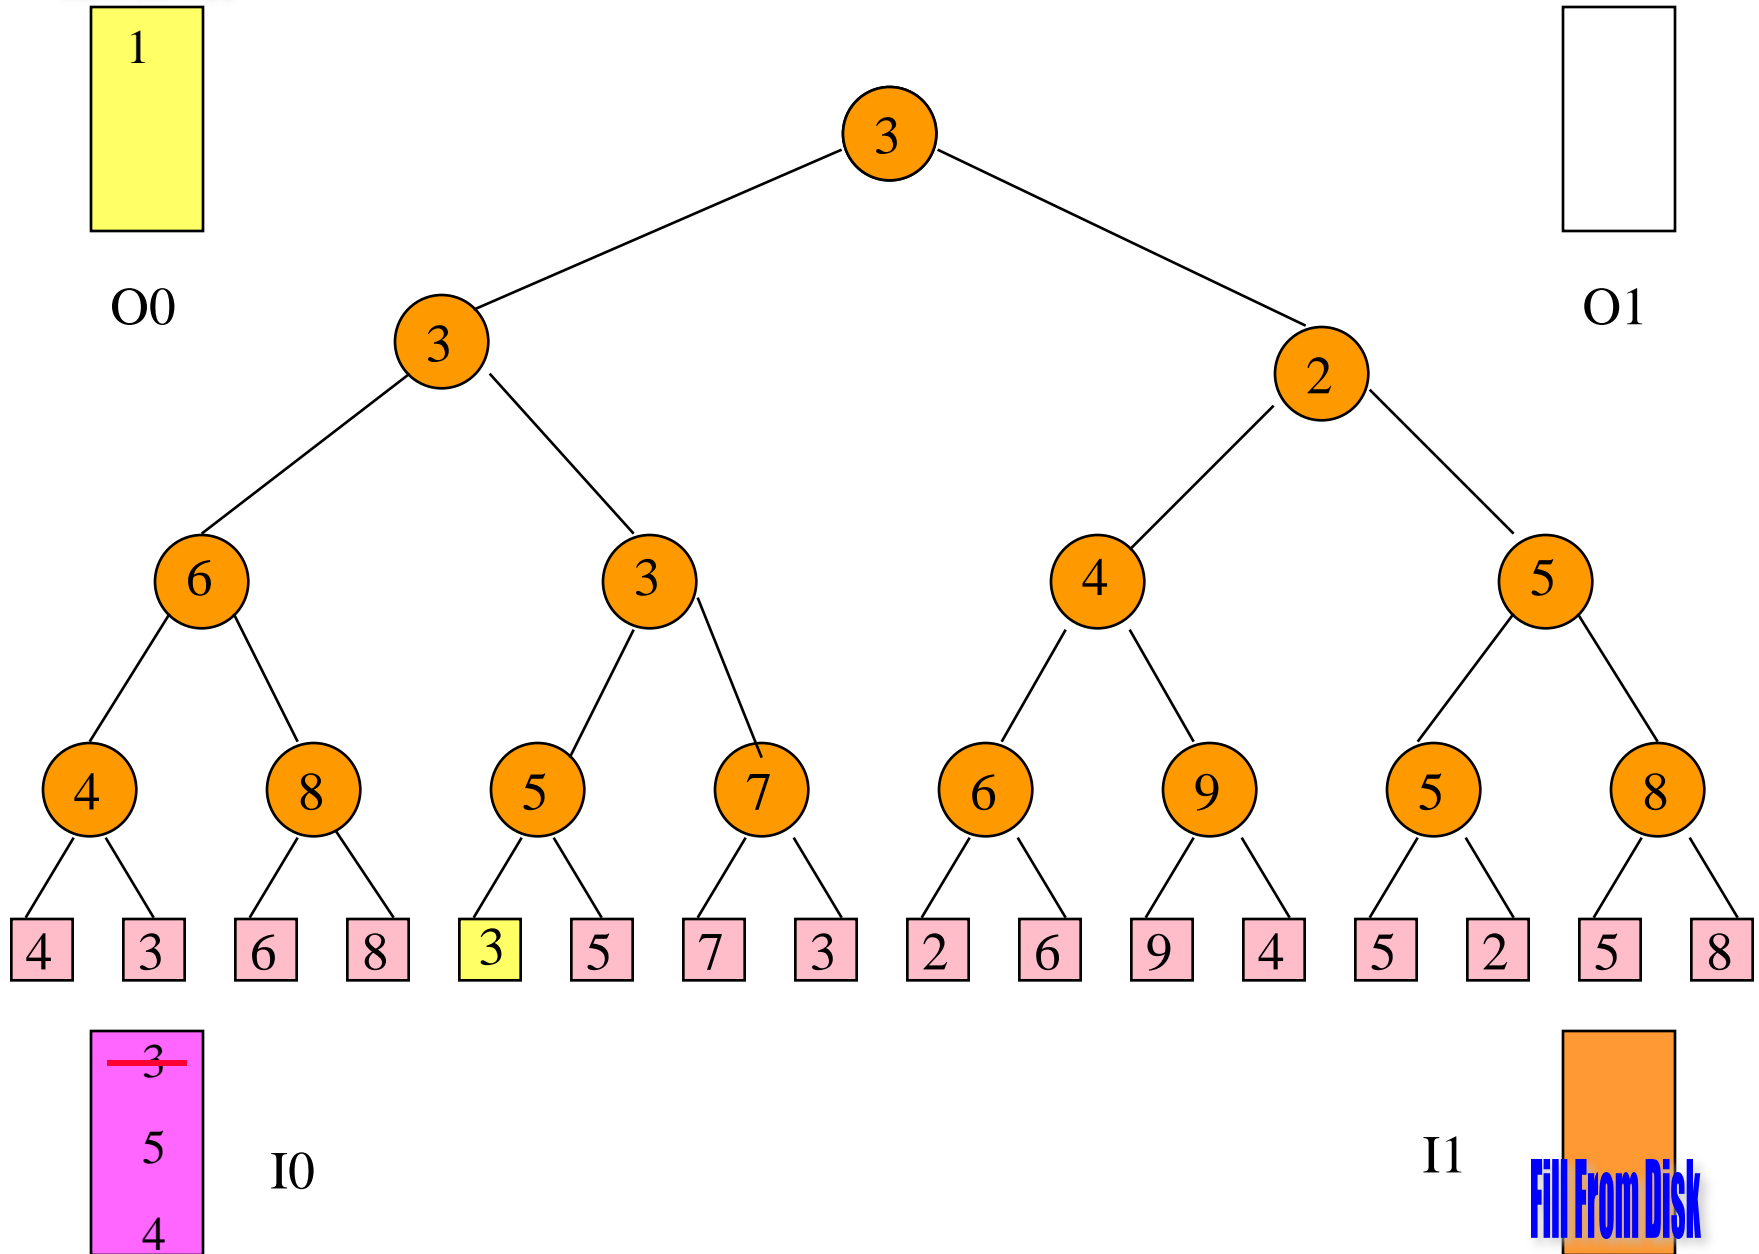
Fill From Tree

# Generate Run 1



Fill From Tree

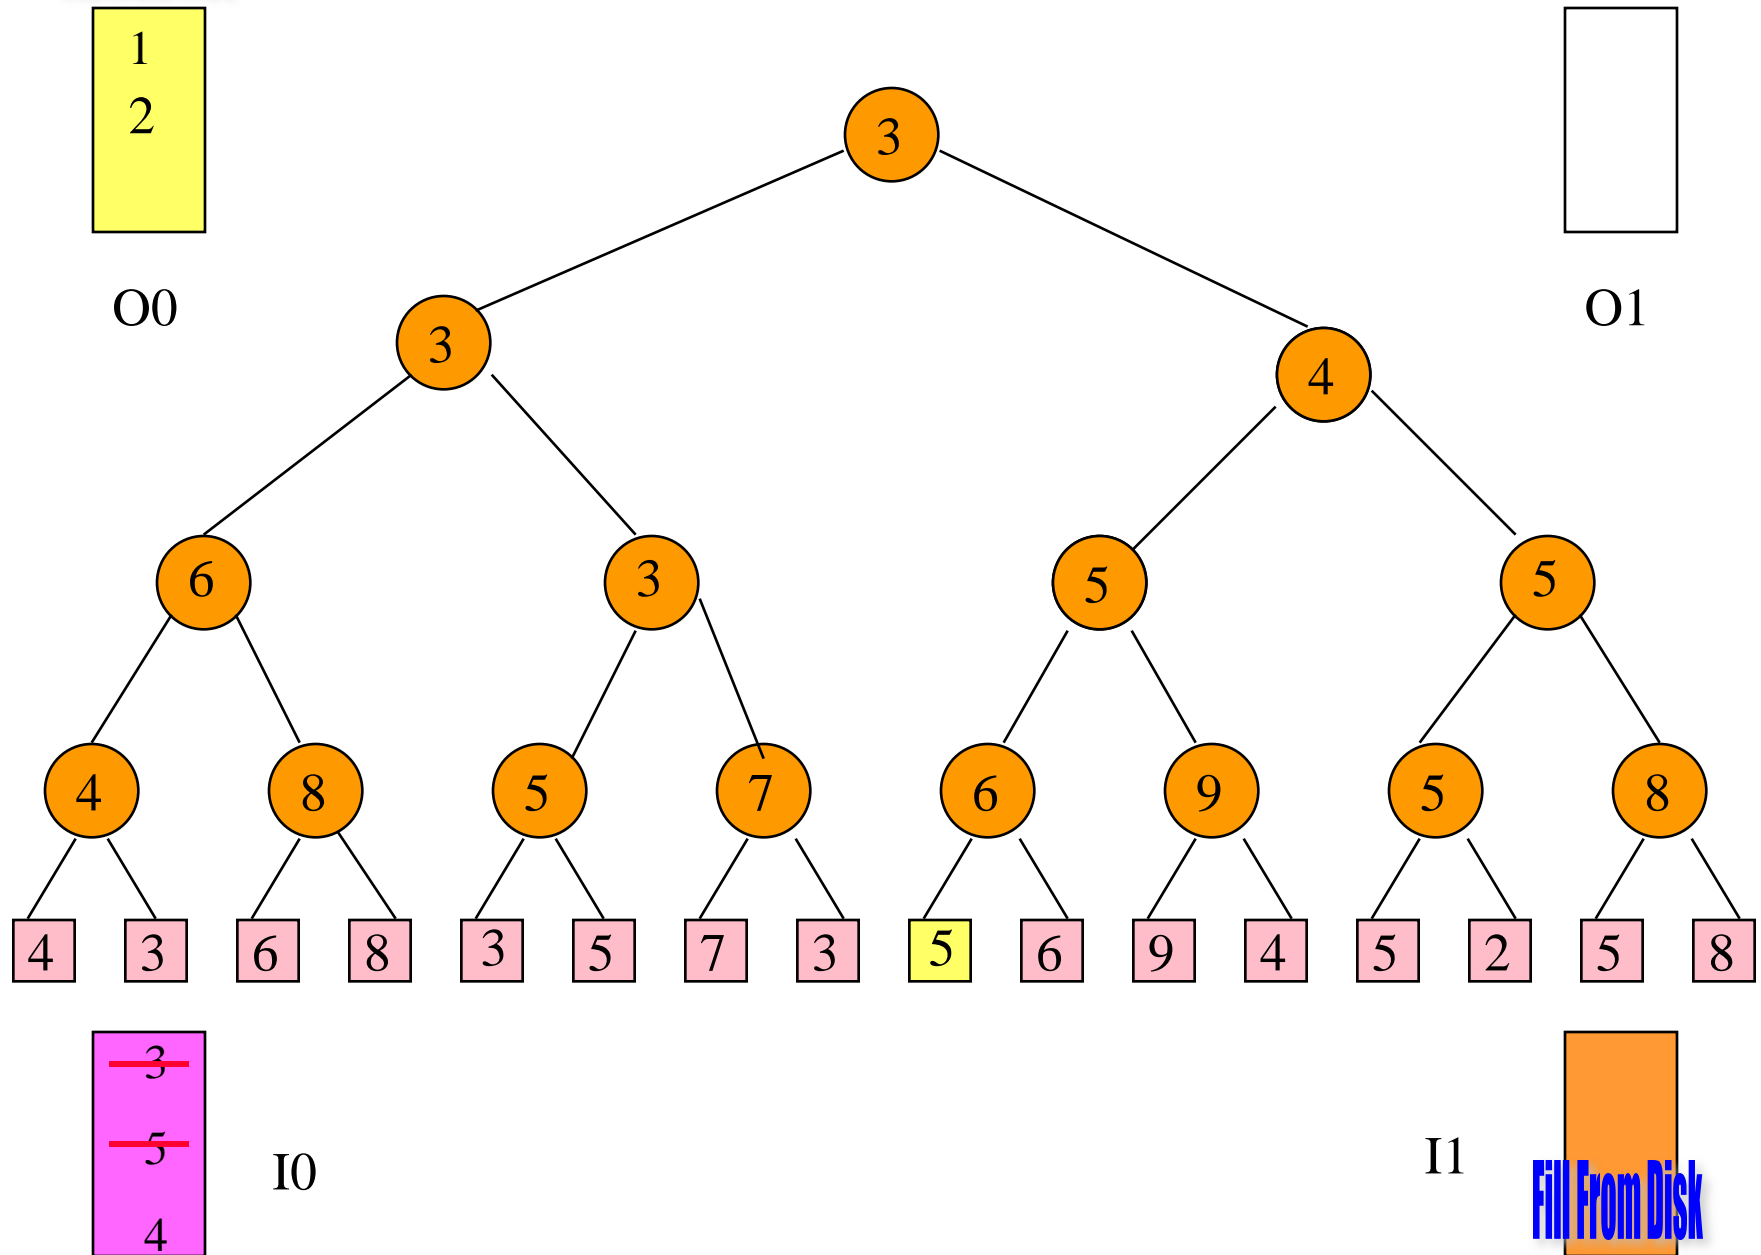
# Generate Run 1






Fill From Tree

# Generate Run 1




Fill From Tree

# Generate Run 1




Write To Disk



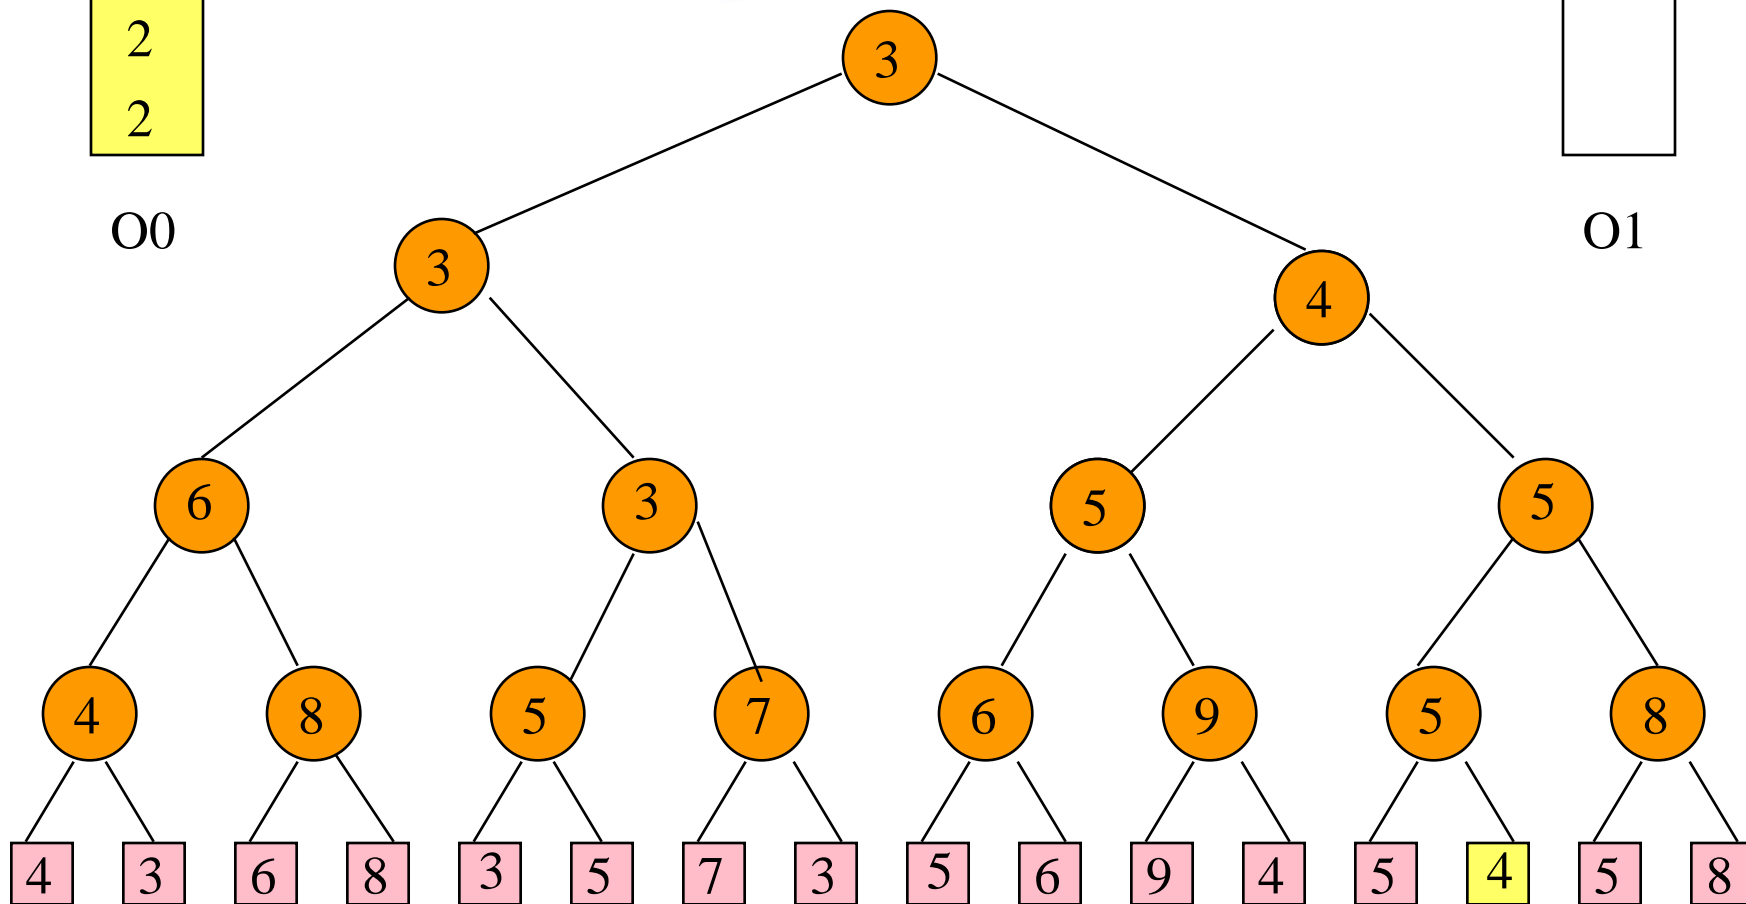
O0

Interchange Role Of Buffers

Fill From Tree



O1



Fill From Disk



I0

I1



Write To Disk

# Continue With Run 1

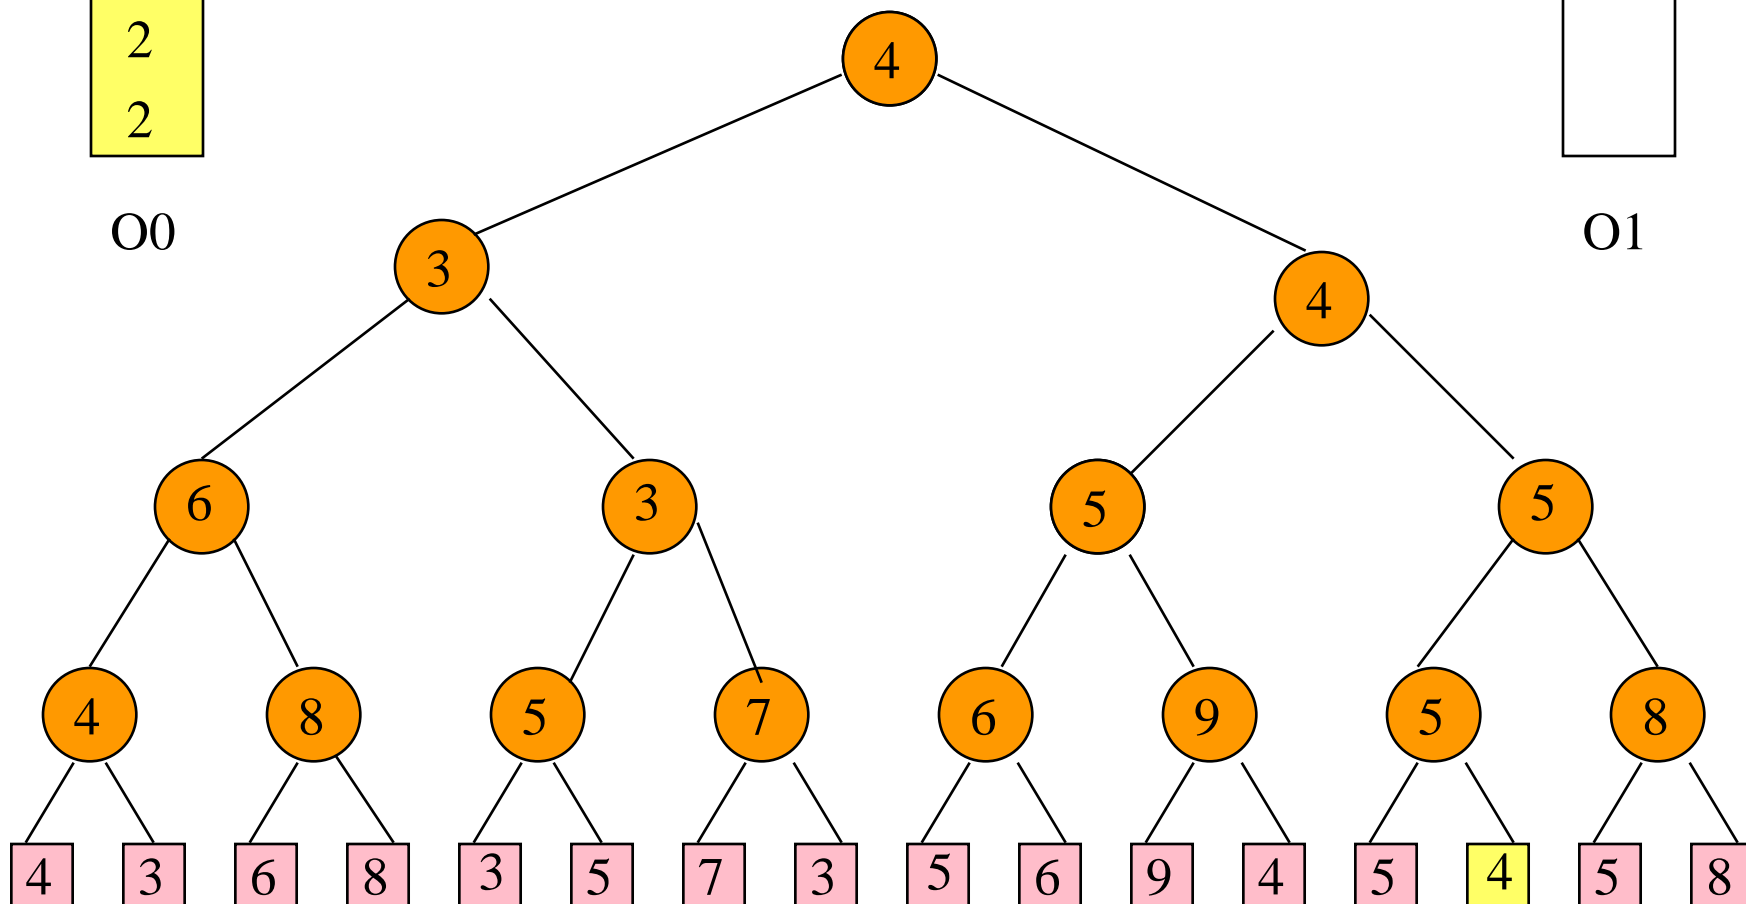
Fill From Tree

1
2
2

--

O0

O1



Fill From Disk

I0

I1

1
9
2

Write To Disk

# Continue With Run 1

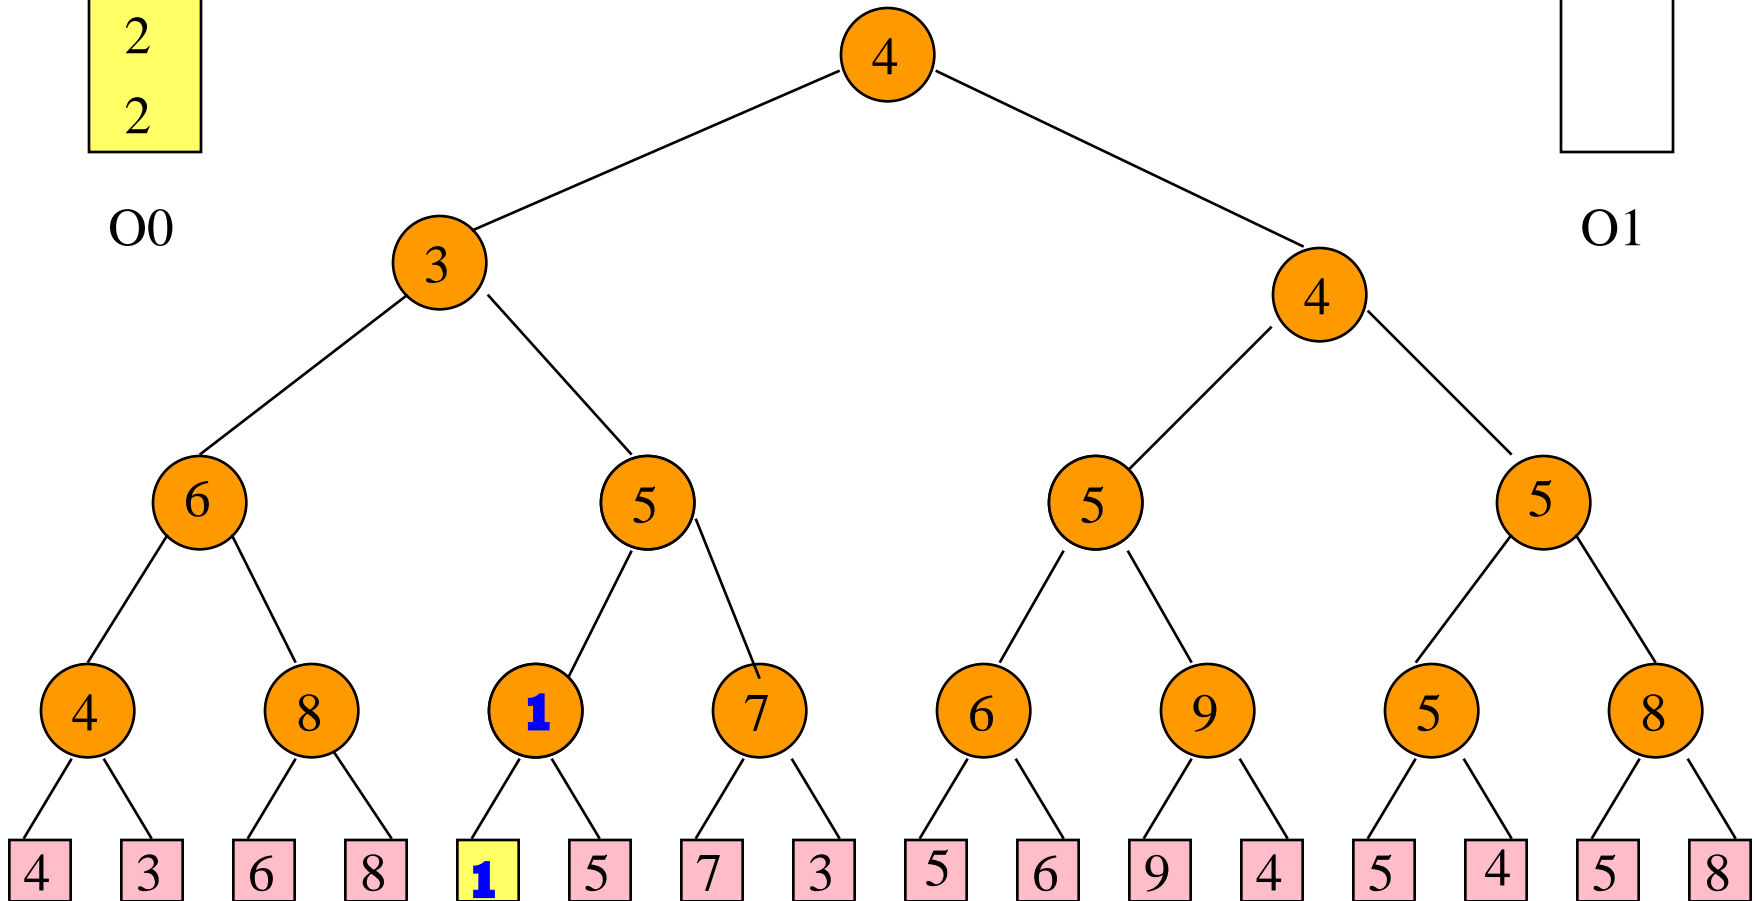
Fill From Tree

1  
2  
2

3

O0

O1



Fill From Disk

I0

I1

~~1~~  
9  
2

Write To Disk

# Continue With Run 1

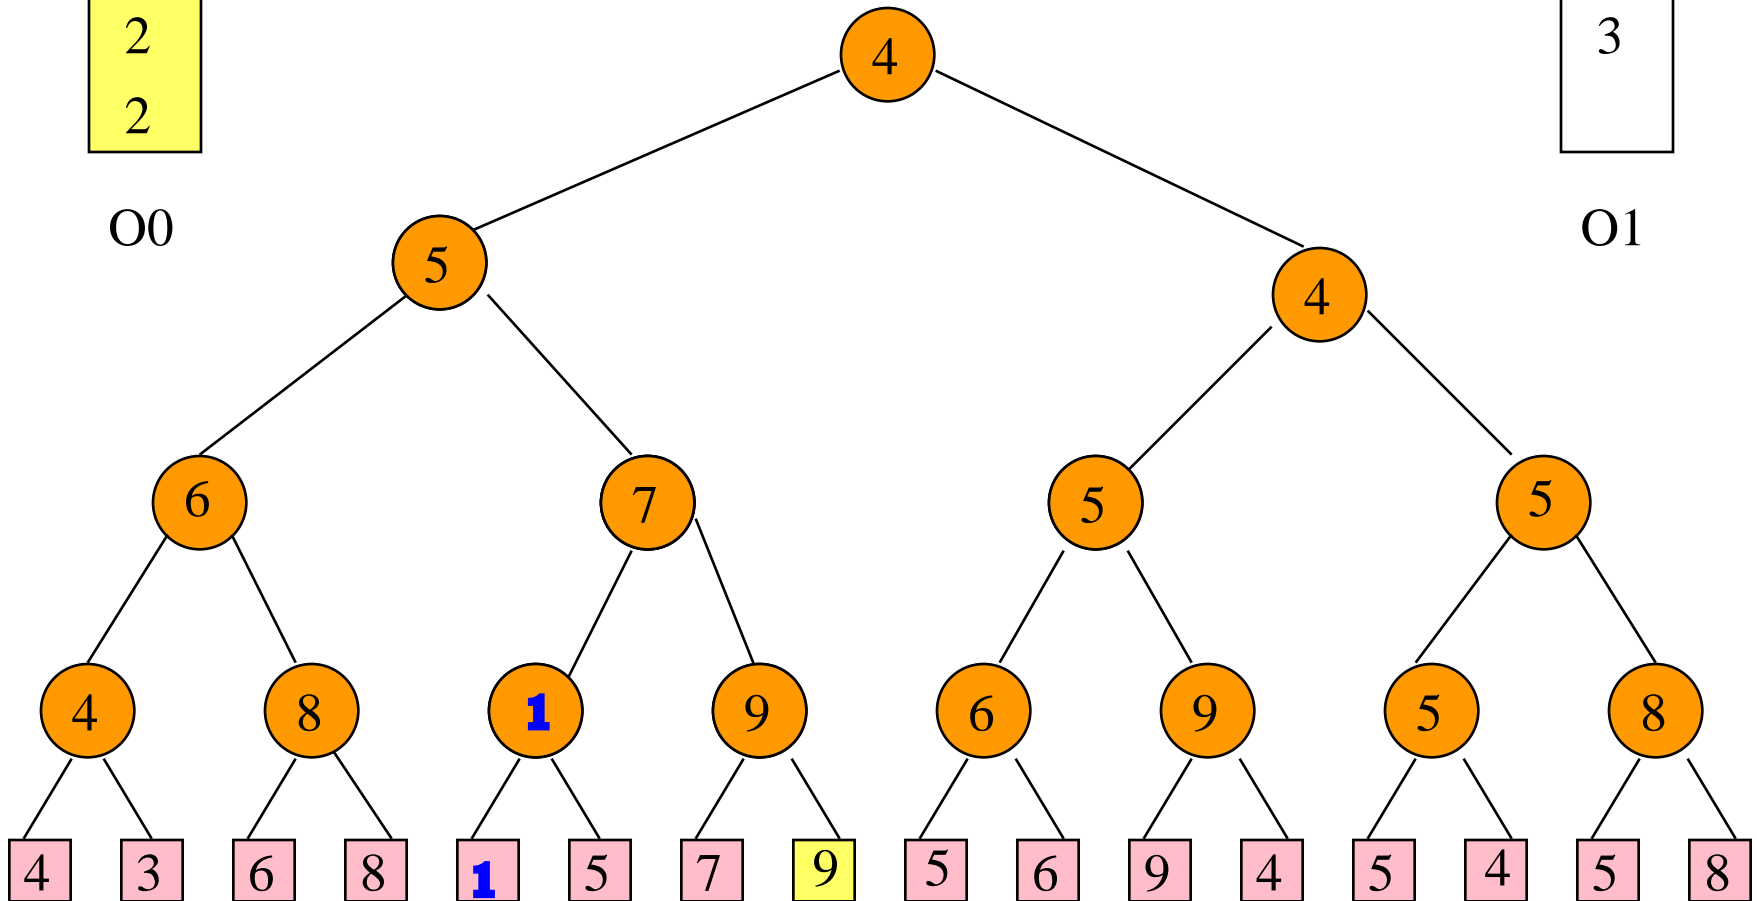
Fill From Tree

1  
2  
2

3  
3

O0

O1



Fill From Disk

I0

I1

~~1~~  
~~9~~  
2

Write To Disk

1  
2  
2

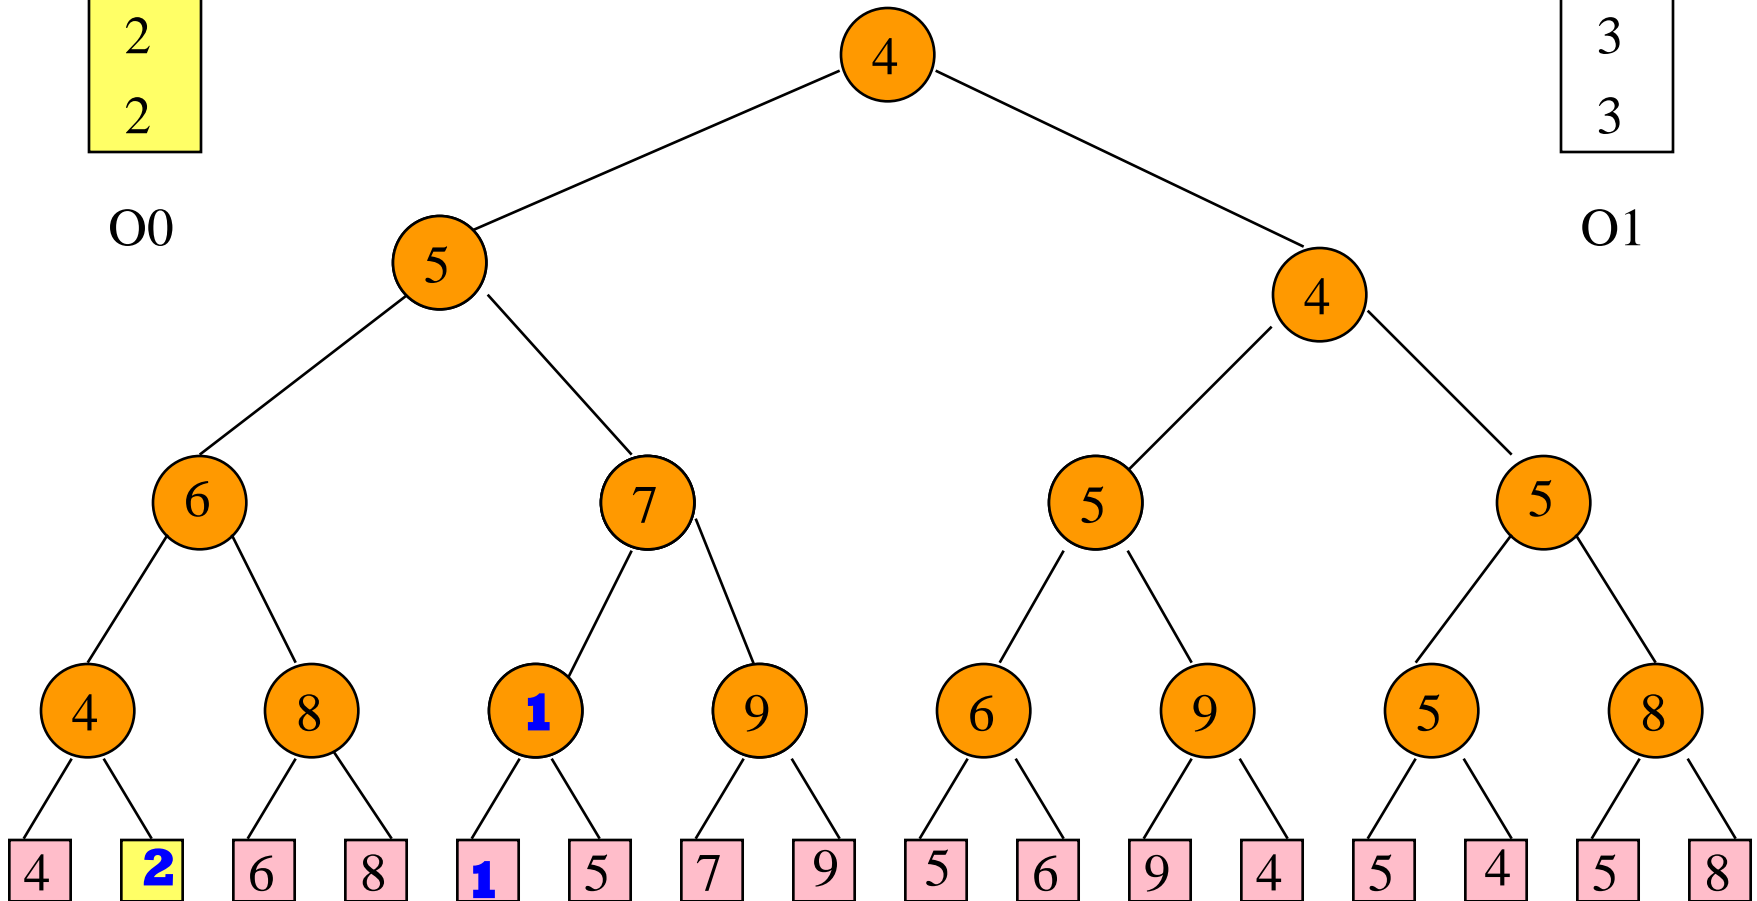
O0

# Continue With Run 1

Fill From Tree

3  
3  
3

O1



Fill From Disk


I0

Interchange Role Of Buffers

I1

~~1~~  
~~9~~  
~~2~~


Fill From Tree




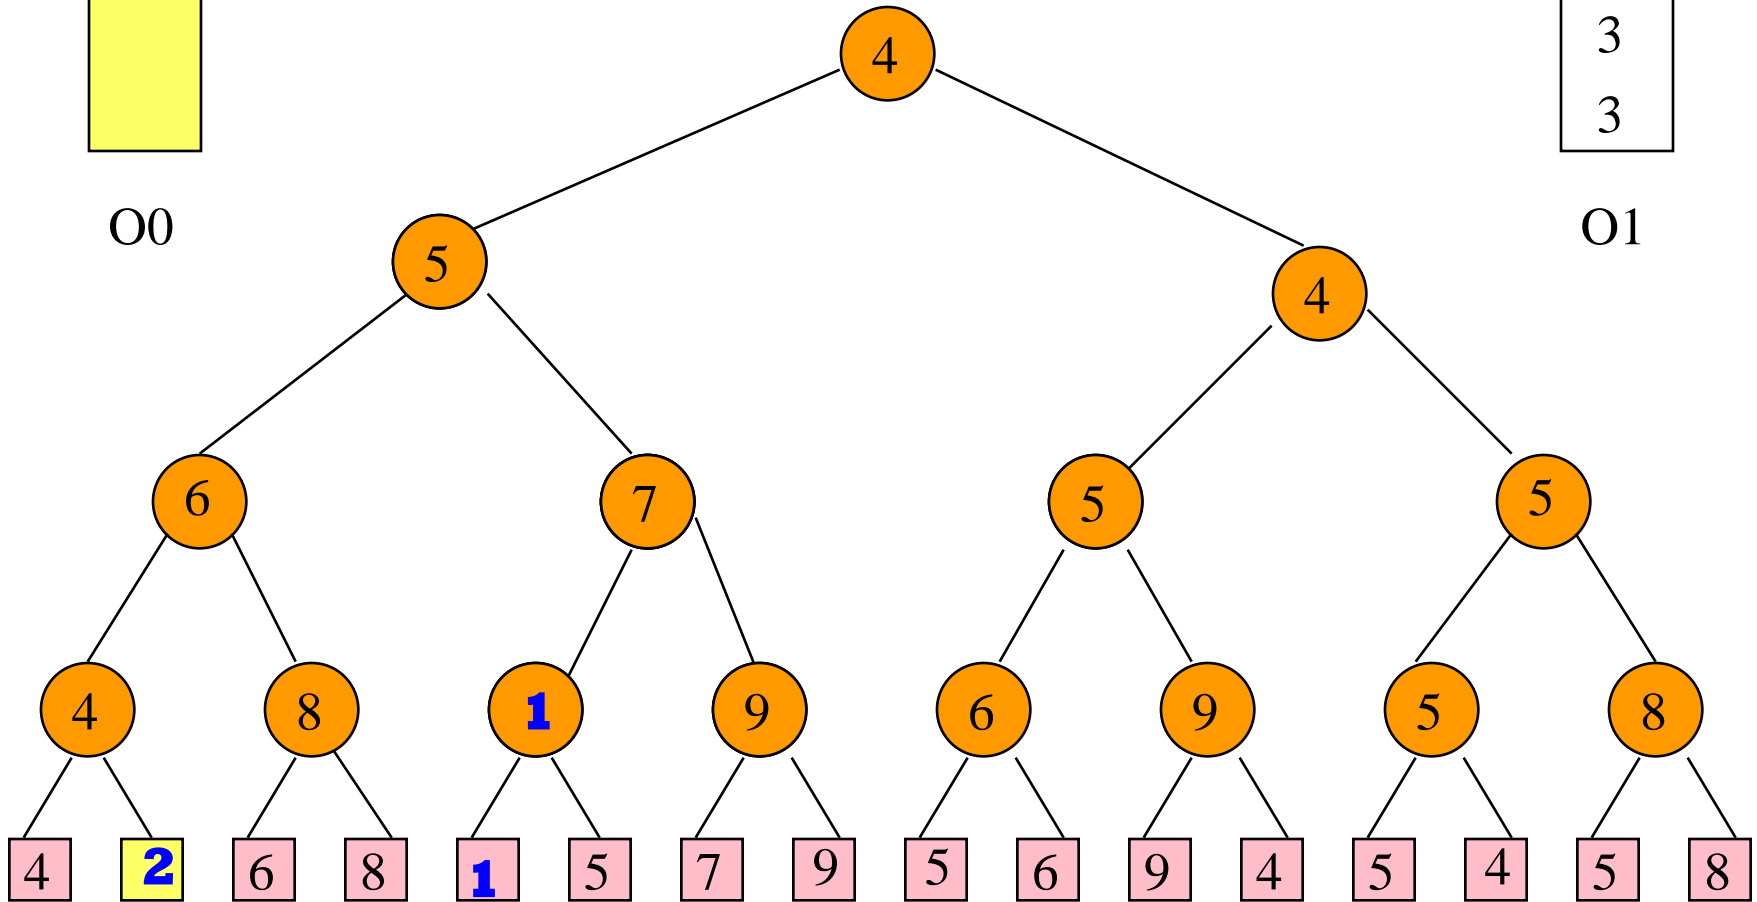
O0

Interchange Role Of Buffers

Write To Disk



O1



I0

I1





Fill From Disk



Fill From Tree


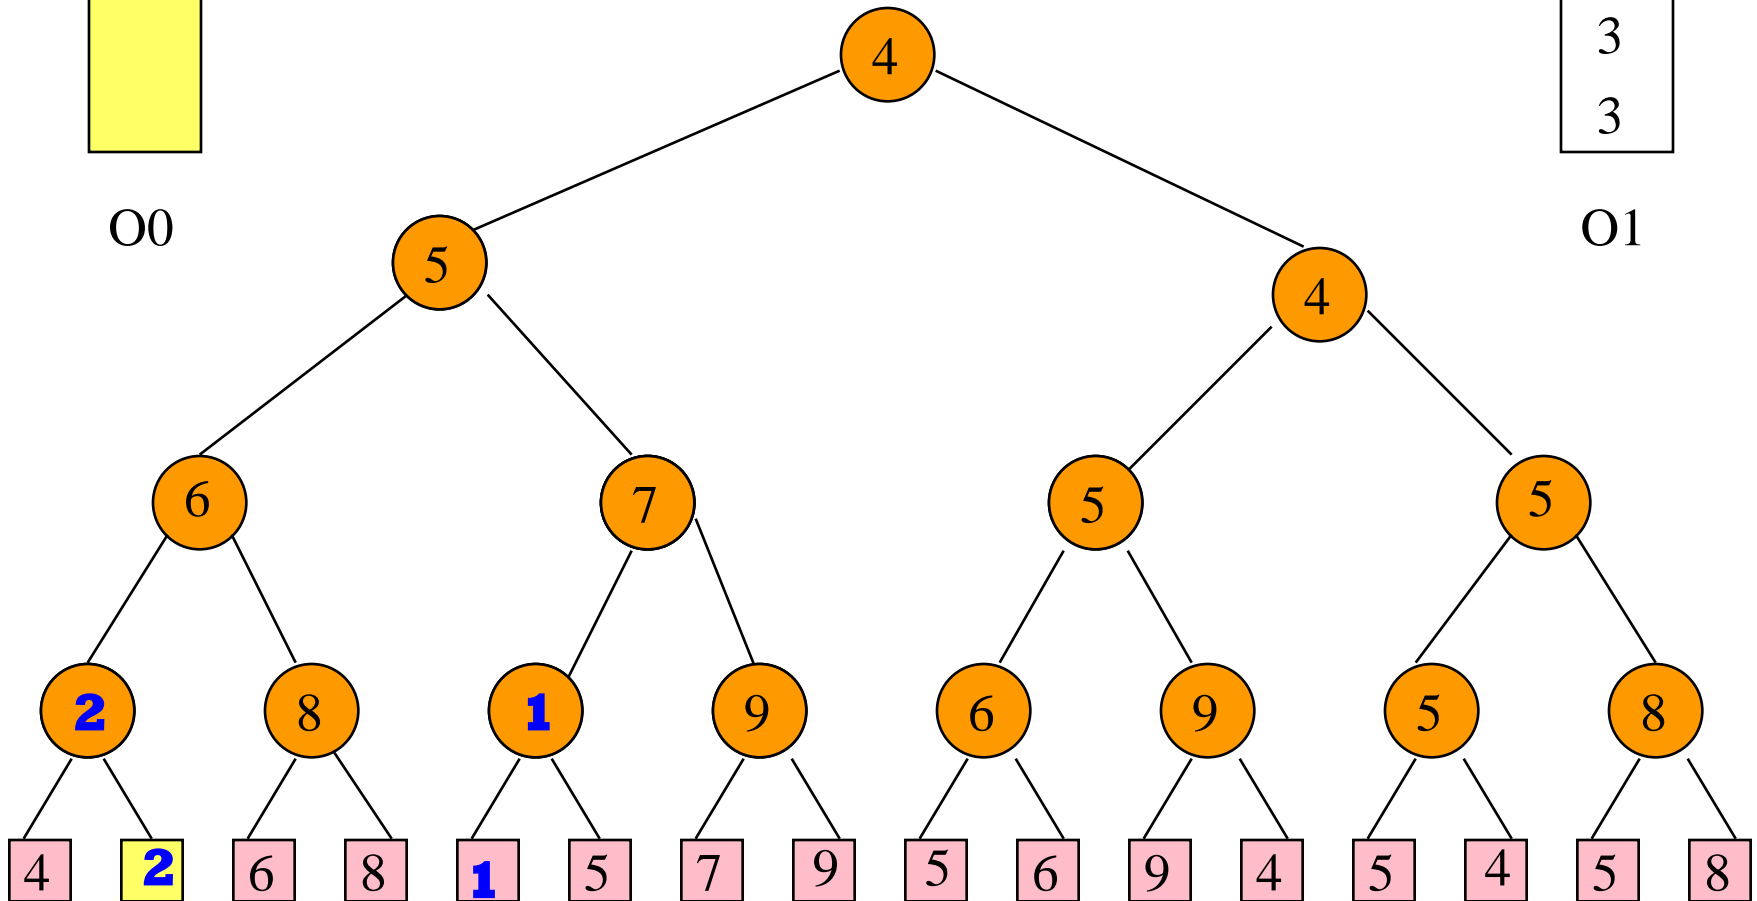
# Continue With Run 1

Write To Disk



O0

O1



I0

I1





Fill From Disk

Fill From Tree


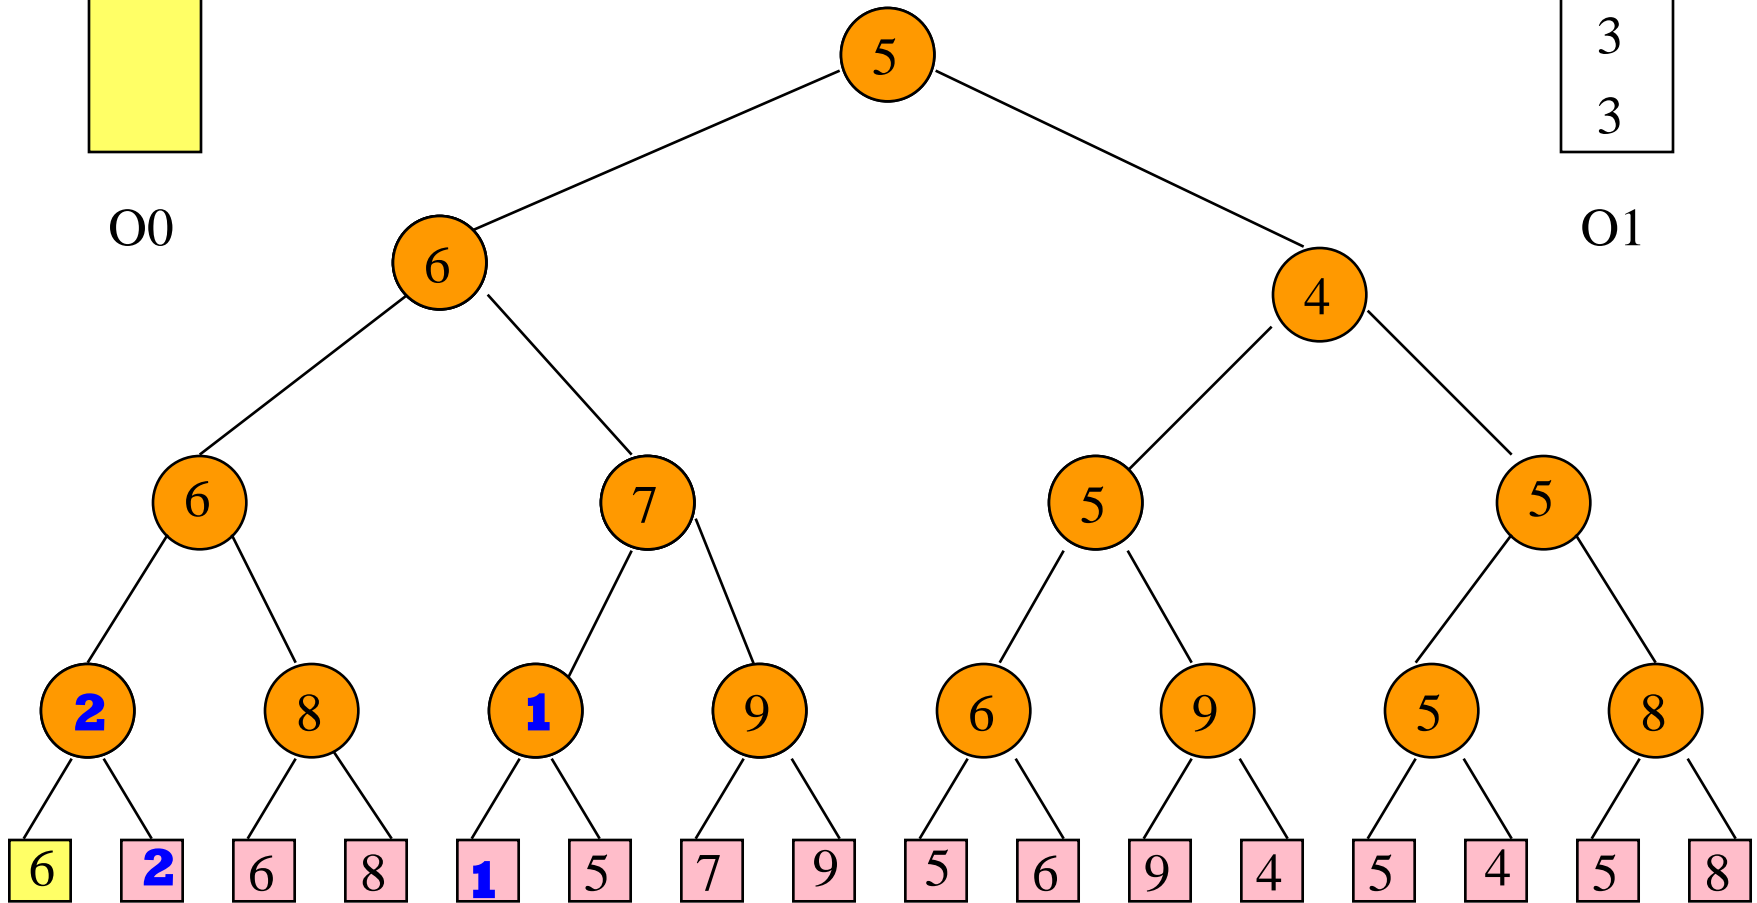
# Continue With Run 1

Write To Disk



O0

O1



I0

I1



Fill From Tree

# Continue With Run 1

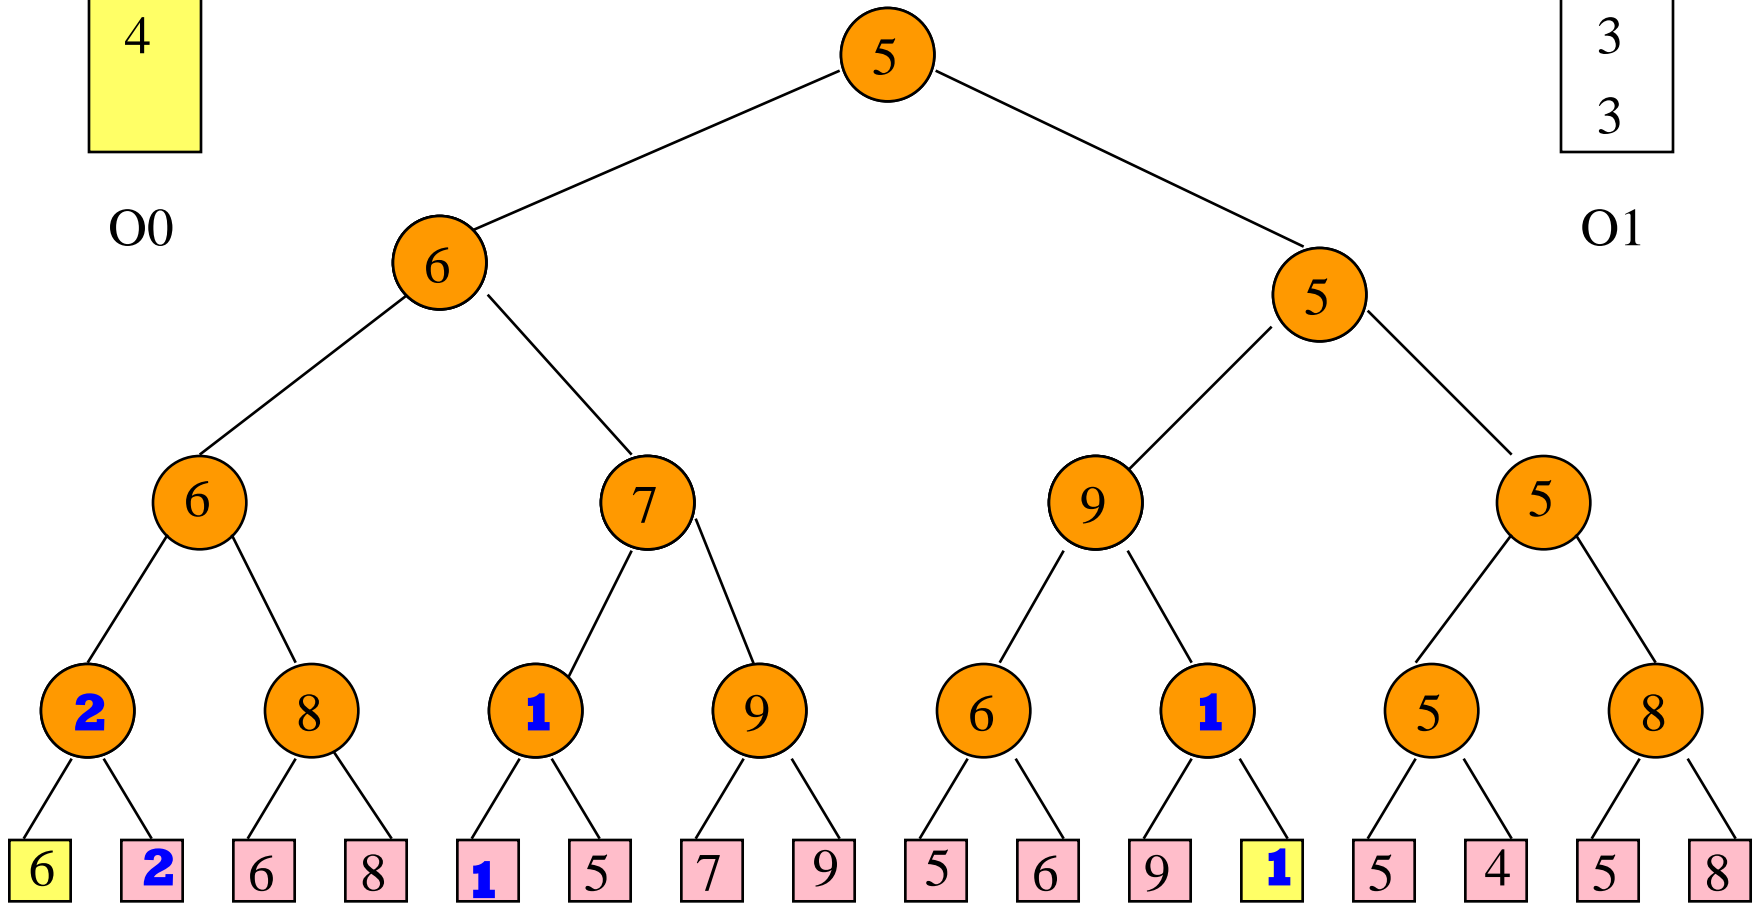
Write To Disk

4
4

3
3
3

O0

O1



<del>6</del>
<del>1</del>
3

I0

I1


--

Fill From Disk


# RUN SIZE

- Let  $k$  be number of external nodes in loser tree.
- Run size  $\geq k$ .
- Sorted input  $\Rightarrow 1$  run.
- Reverse of sorted input  $\Rightarrow n/k$  runs.
- Average run size is  $\sim 2k$ .

# Merging Runs Of Different Length



Cost = 44



Cost = 42

**Best merge sequence?**

# Requirements

- 3 buffers for improving run generation
- Run lengths and best merge sequence should be output
- Performance comparison
- `#include <thread>`

# Phase 5

Merge Sort: Improve Run Merging


# Improve Run Merging

- Reduce number of merge passes.
  - Use higher order merge.
  - Number of passes  
=  $\text{ceil}(\log_k(\text{number of initial runs}))$   
where  $k$  is the merge order.
- More generally, a higher-order merge reduces the cost of the optimal merge tree.

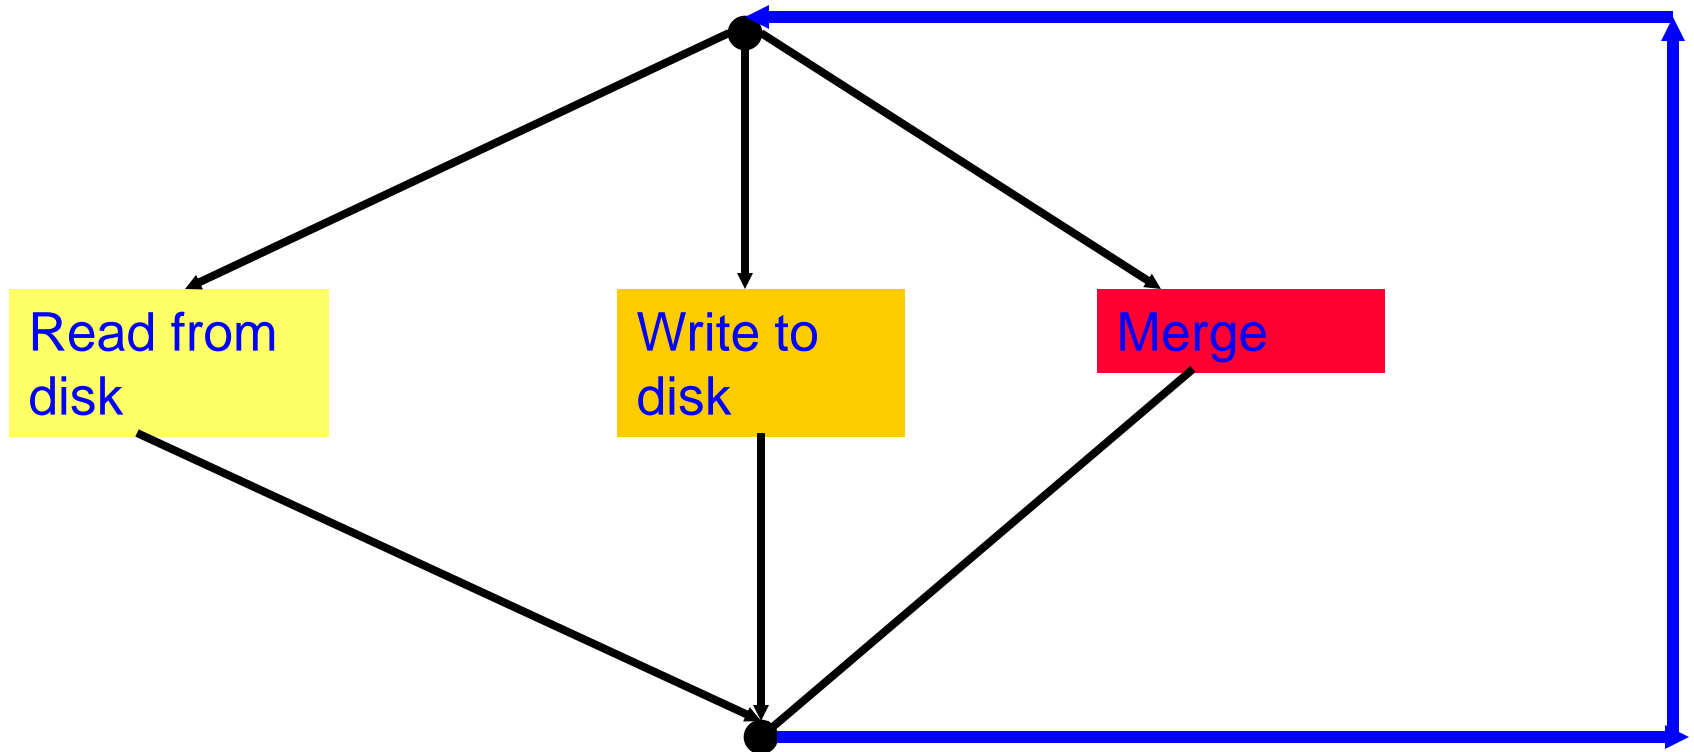


# Improve Run Merging


- Overlap input, output, and internal merging.



# Steady State Operation



# Partitioning Of Memory



- Need exactly **2** output buffers.
- Need at least  **$k+1$**  ( **$k$**  is merge order) input buffers.
- **$2k$**  input buffers suffice.


# Number Of Input Buffers

- When  $2$  input buffers are dedicated to each of the  $k$  runs being merged,  $2k$  buffers are not enough!
- Input buffers must be allocated to runs on an *as needed* basis.


# Buffer Allocation

- When ready to read a buffer load, determine which run will exhaust first.
  - Examine key of the last record read from each of the  $k$  runs.
  - Run with smallest last key read will exhaust first.
- Next buffer load of input is to come from run that will exhaust first, allocate an input buffer to this run.


# Buffer Layout



Output  
buffers



Input buffer  
queues  $k=9$



Pool of free input buffers

# Initialize To Merge $k$ Runs

- Initialize  $k$  queues of input buffers, 1 queue per run, 1 buffer per run.
- Input one buffer load from each of the  $k$  runs.
- Put  $k - 1$  unused input buffers into pool of free buffers.
- Set  $\text{activeOutputBuffer} = 0$ .
- Initiate input of next buffer load from first run to exhaust. Use remaining unused input buffer for this input.

# The Method kWayMerge

- k-way merge from input queues to the active output buffer.
- Merge stops when either the output buffer gets full or when an end-of-run key is merged into the output buffer.
- If merge hasn't stopped and an input buffer gets empty, advance to next buffer in queue and free empty buffer.



# Merge k Runs

repeat

kWayMerge;

wait for input/output to complete;

add new input buffer (if any) to queue for its run;

determine run that will exhaust first;

if (there is more input from this run)

initiate read of next block for this run;

initiate write of active output buffer;

activeOutputBuffer = 1 – activeOutputBuffer;

until end-of-run key merged;

# What Can Go Wrong?

## kWayMerge

- k-way merge from input queues to the active output buffer.
- Merge stops when either the output buffer gets full or when an end-of-run key is merged into the output buffer.
- If merge hasn't stopped and an input buffer gets empty, advance to next buffer in queue and free empty buffer.  
There may be no next buffer in the queue.

# What Can Go Wrong?

## Merge k Runs

repeat

kWayMerge;

wait for input/output to complete;

add new input buffer (if any) to queue for its run;

determine run that will exhaust first;

if (there is more input from this run)

initiate read of next block for this run;

initiate write of active output buffer;

$\text{activeOutputBuffer} = 1 - \text{activeOutputBuffer};$

until end of run key merged;

There may be  
no free input  
buffer to read  
into.

# kWayMerge

- If merge hasn't stopped and an input buffer gets empty, advance to next buffer in queue and free empty buffer.  
There may be no next buffer in the queue.
- If this type of failure were to happen, using two different and valid analyses, we will end up with inconsistent counts of the amount of data available to kWayMerge.
- Data available to kWayMerge is data in
  - Input buffer queues.
  - Active output buffer.
  - Excludes data in buffer being read or written.

# No Next Buffer In Queue

repeat

kWayMerge;



wait for input/output to complete;

add new input buffer (if any) to queue for its run;

determine run that will exhaust first;

if (there is more input from this run)

initiate read of next block for this run;

initiate write of active output buffer;

activeOutputBuffer = 1 – activeOutputBuffer;

until end-of-run key merged;

- Exactly **k** buffer loads available to kWayMerge.

# kWayMerge

- If merge hasn't stopped and an input buffer gets empty, advance to next buffer in queue and free empty buffer.  
There may be no next buffer in the queue.
- Alternative analysis of data available to kWayMerge at time of failure.
  - $< 1$  buffer load in active output buffer
  - $\leq k - 1$  buffer loads in remaining  $k - 1$  queues
  - Total data available to k-way merge is  $< k$  buffer loads.

# Merge $k$ Runs

initiate read of next block for this run;

There may be  
no free input  
buffer to read  
into.

- Suppose there is no free input buffer.
- One analysis will show there are exactly  $k + 1$  buffer loads in memory (including newly read input buffer) at time of failure.
- Another analysis will show there are  $> k + 1$  buffer loads in memory at time of failure.
- Note that at time of failure there is no buffer being read or written.

# No Free Input Buffer

repeat

kWayMerge;

wait for input/output to complete;

add new input buffer (if any) to queue for its run;

determine run that will exhaust first;

if (there is more input from this run)

initiate read of next block for this run;



initiate write of active output buffer;

activeOutputBuffer = 1 – activeOutputBuffer;

until end-of-run key merged;

- Exactly  $k + 1$  buffer loads in memory.



# Merge $k$ Runs

initiate read of next block for this run;

There may be no free input buffer to read into.

- Alternative analysis of data in memory.
  - 1 buffer load in the active output buffer.
  - 1 input queue may have an empty first buffer.
  - Remaining  $k - 1$  input queues have a nonempty first buffer.
  - Remaining  $k$  input buffers must be in queues and full.
  - Since  $k > 1$ , total data in memory is  $> k + 1$  buffer loads.

# Minimize Wait Time For I/O To Complete

Time to fill an output buffer

~ time to read a buffer load

~ time to write a buffer load

# Initializing For Next k-way Merge

Change

**if** (there is more input from this run)

    initiate read of next block for this run;

to

**if** (there is more input from this run)

    initiate read of next block for this run;

**else**

    initiate read of a block for the next k-way merge;

# Requirements

- Allocate as needed strategy
- Performance comparison & analysis